

Deep Reinforcement Learning for Discrete Action Space

Group 6

Wei Li Ming-Xu Huang

Reference

- Mnih, Volodymyr, et al. "**Human-level control through deep reinforcement learning.**" Nature 518.7540 (2015): 529-533.
- Van Hasselt, Hado, Arthur Guez, and David Silver. "**Deep Reinforcement Learning with Double Q-Learning.**" AAAI. 2016.
- Wang, Ziyu, et al. "**Dueling network architectures for deep reinforcement learning.**" arXiv preprint arXiv:1511.06581 (2016).
- All the papers are published by **Google DeepMind**

Outline

- Introduction
- Deep Q Network
- Double DQN
- Dueling Network
- Experiment
- Conclusion

Outline

- Introduction
- Deep Q Network
- Double DQN
- Dueling Network
- Experiment
- Conclusion

Introduction

- In this slide we will discuss the traditional DQN and various improvements to Deep Q Network
- We will discuss the results of each methods and compare their performance.

Outline

- Introduction
- Deep Q Network
 - Target-Q
 - Experience replay
- Double DQN
- Dueling Network
- Experiment
- Conclusion

Deep Q Network

- State value: $V(s)$
- Action value: $Q(s, a)$
- Approximate action value by a neural network parameterized by θ
 - $Q(s, a; \theta)$
- Objective function
 - $L(\theta) = \mathbb{E}_{s,a,r,s'} \left[\left(y^{DQN} - Q(s, a; \theta) \right)^2 \right]$
 - $y^{DQN} = r + \gamma \max_{a'} Q(s', a'; \theta)$
- Gradient
 - $\nabla L(\theta) = \mathbb{E}_{s,a,r,s'} [(y^{DQN} - Q(s, a; \theta)) \nabla Q(s, a; \theta)]$

Deep Q Network(cont.)

- Target-Q Network
 - Small updates to Q
 - Significantly change the policy
 - Changing correlations between the action-values and the target values
 - Neural networks is to use a separate network for generating the targets y
 - Every C step, clone weights θ of behavior Q network to target Q network weights θ^-

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

$$\text{Set } y_j = \begin{cases} r_j \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) \end{cases}$$

Every C steps reset $\hat{Q} = Q$

Deep Q Network(cont.)

- Experience replay
 - Store experiences e_t in a fixed size buffer D
 - $e_t = (s_t, a_t, r_t, s_{t+1})$
 - $D = \{e_1, e_2, \dots, e_n\}$
 - Trained by randomly sampling mini-batch of experiences from buffer uniformly
 - Decreasing the correlations present in the sequence of observations
 - Updating at iteration i uses the following loss function

- $$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Outline

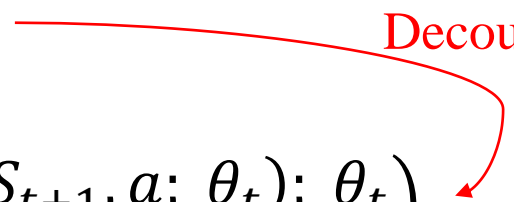
- Introduction
- Deep Q Network
- Double DQN
 - Double Q-Learning
 - Overestimation of Q-Learning
 - Double DQN
- Dueling Network
- Experiment
- Conclusion

Double DQN

- The **max operator** in standard Q-learning and DQN, uses the same values both to select and to evaluate an action.
- This makes it **more likely to select overestimated values**, resulting in overoptimistic value estimates.
- To prevent this, we can decouple the selection from the evaluation.

Q-learning: $Y_t^Q \equiv R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t)$

$Y_t^Q \equiv R_{t+1} + \gamma Q(S_{t+1}, \operatorname{argmax}_a Q(S_{t+1}, a; \theta_t); \theta_t)$




Double Q-learning: $Y_t^{DoubleQ} \equiv R_{t+1} + \gamma Q(S_{t+1}, \operatorname{argmax}_a Q(S_{t+1}, a; \theta_t); \theta'_t)$

Double DQN

- Overestimation of Q-Learning
 - Consider a real-valued continuous state space with 10 discrete actions in each state.
 - For simplicity, the true optimal action values in this example depend only on state so that **in each state all actions have the same true value**.

Double DQN(cont.)

True value and an estimate



$$Q_*(s, a) = \sin(s)$$

Degree of polynomial: 6

Different True Value Function



$$Q_*(s, a) = 2 \exp(-s^2)$$


Degree of polynomial: 6

Double DQN(cont.)



$$Q_*(s, a) = 2 \exp(-s^2)$$

Degree of polynomial: 6

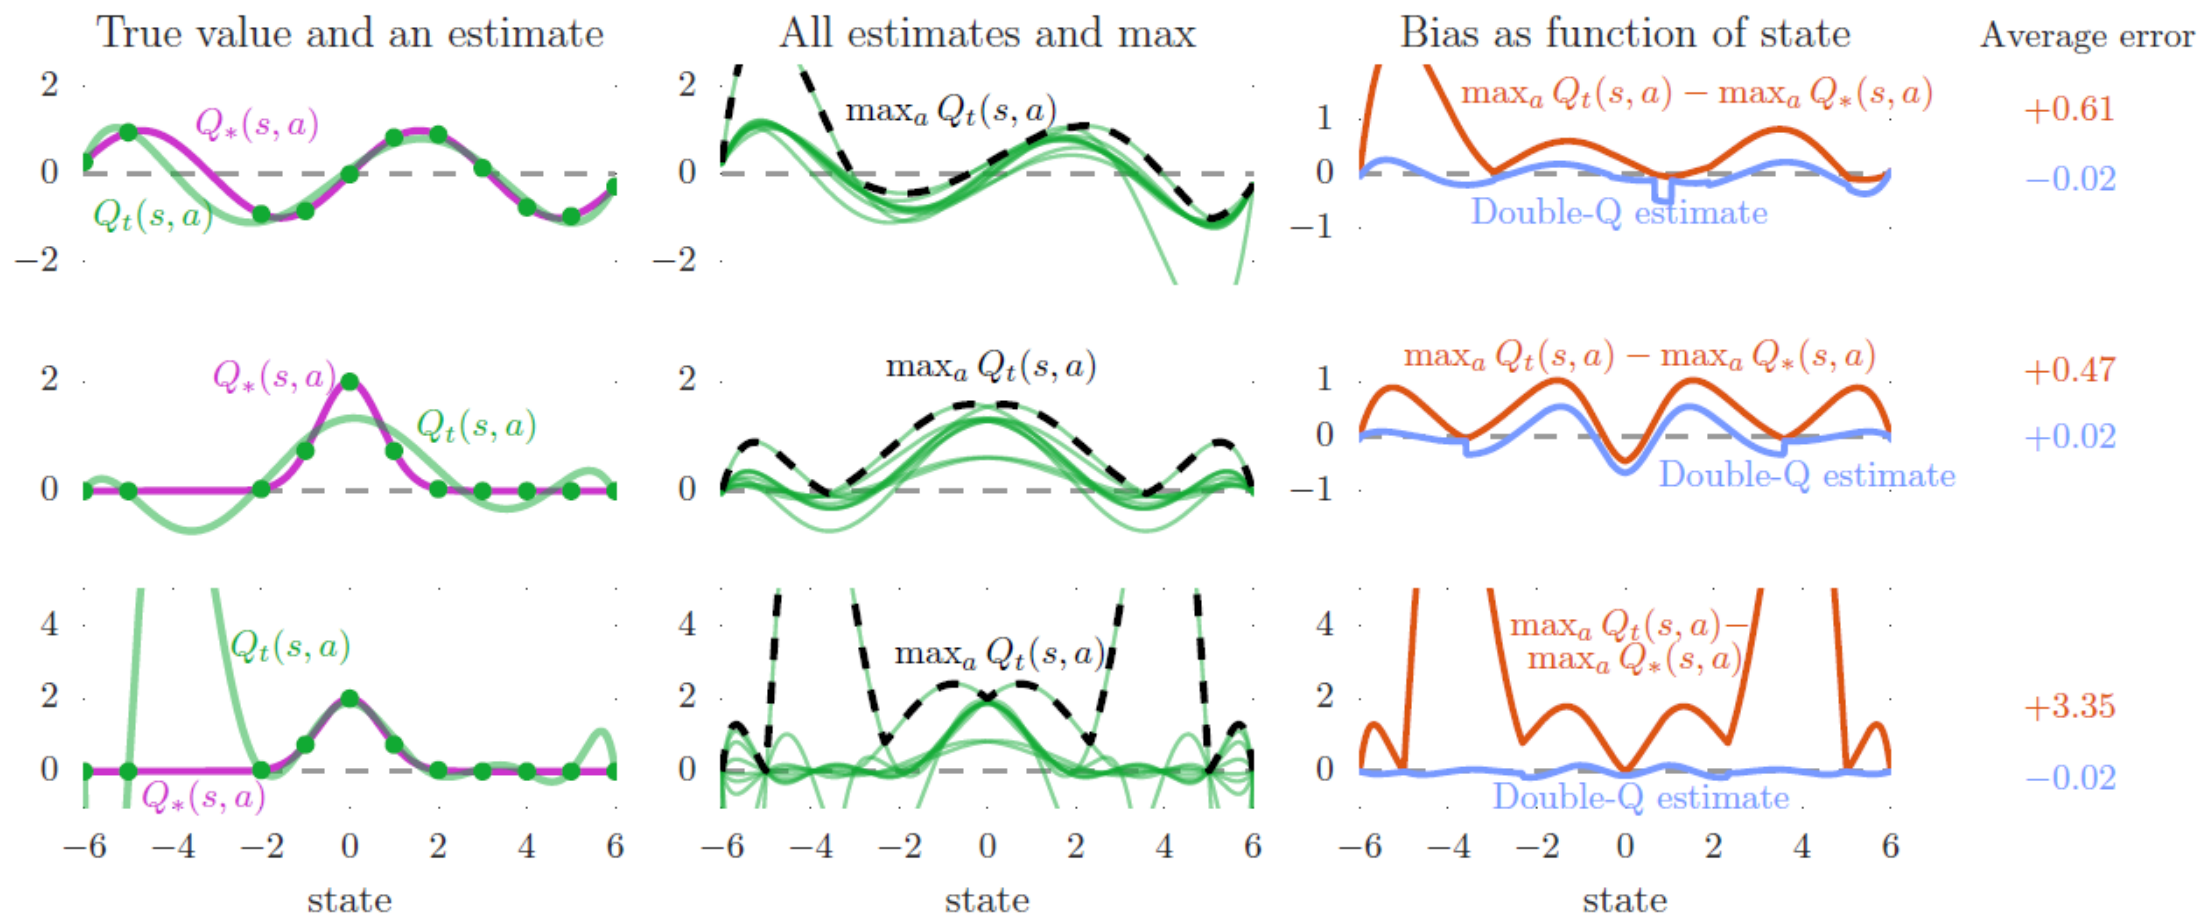


$$Q_*(s, a) = 2 \exp(-s^2)$$

Degree of polynomial: 9

Different Degree of Polynomial

Double DQN(cont.)




Double DQN(cont.)

- The idea of Double Q-learning is to reduce overestimations by decomposing the max operation in the target into action selection and action evaluation.
- Evaluate the greedy policy according to the online network.
- Using the target network to estimate its value.

$$Y_t^{DQN} \equiv R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t^-)$$

$$Y_t^{DQN} \equiv R_{t+1} + \gamma Q\left(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t^-); \theta_t^-\right)$$

$$Y_t^{DoubleDQN} \equiv R_{t+1} + \gamma Q\left(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t); \theta_t^-\right)$$


Double DQN(cont.)

Algorithm 1: Double DQN Algorithm.

input : \mathcal{D} – empty replay buffer; θ – initial network parameters, θ^- – copy of θ

input : N_r – replay buffer maximum size; N_b – training batch size; N^- – target network replacement freq.

for episode $e \in \{1, 2, \dots, M\}$ **do**

 Initialize frame sequence $\mathbf{x} \leftarrow ()$

for $t \in \{0, 1, \dots\}$ **do**

 Set state $s \leftarrow \mathbf{x}$, sample action $a \sim \pi_{\mathcal{B}}$

 Sample next frame x^t from environment \mathcal{E} given (s, a) and receive reward r , and append x^t to \mathbf{x}

if $|\mathbf{x}| > N_f$ **then** delete oldest frame $x_{t_{min}}$ from \mathbf{x} **end**

 Set $s' \leftarrow \mathbf{x}$, and add transition tuple (s, a, r, s') to \mathcal{D} ,

 replacing the oldest tuple if $|\mathcal{D}| \geq N_r$

 Sample a minibatch of N_b tuples $(s, a, r, s') \sim \text{Unif}(\mathcal{D})$

 Construct target values, one for each of the N_b tuples:

 Define $a^{\max}(s'; \theta) = \arg \max_{a'} Q(s', a'; \theta)$

$y_j = \begin{cases} r & \text{if } s' \text{ is terminal} \\ r + \gamma Q(s', a^{\max}(s'; \theta); \theta^-) & \text{otherwise.} \end{cases}$

 Do a gradient descent step with loss $\|y_j - Q(s, a; \theta)\|^2$

 Replace target parameters $\theta^- \leftarrow \theta$ every N^- steps

end



end

Outline

- Introduction
- Deep Q Network
- Double DQN
- Dueling Network
 - Advantage Function
 - Dueling network architecture
 - Combine methods
- Experiment
- Conclusion


Dueling Network

- Advantage Function
 - $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$
- The value stream learns to pay attention to the road.
- The advantage stream learns to pay attention only when there are cars immediately in front, so as to avoid collisions



Dueling Network(cont.)

- Produce a state value and advantage functions via a single network
 - state value is a scalar
 - advantage functions is a vector of size $|A|$
- Combine state value and advantage functions to generate action values
 - $Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \alpha) + A(s, a; \theta, \beta)$



Dueling Network(cont.)

- Unidentifiable
 - $Q(s, a) = V(s) + A(s, a)$
 - However, $V(s) + A(s, a) = (V(s) - C) + (A(s, a) + C)$ (Poor performance)
 - Improvement
 - $Q(s, a) = V(s) + \left(A(s, a) - \max_{a'} A(s, a') \right)$
 - $a^* = \max_{a'} Q(s, a'), Q(s, a^*) = V(s)$
 - $Q(s, a) = V(s) + \left(A(s, a) - \frac{1}{|A|} \sum_{a'} A(s, a') \right)$
 - $a^* = \max_{a'} Q(s, a'), Q(s, a^*) \neq V(s)$
 - Increase stability of the optimization

Dueling Network(cont.)

- Strengths
 - Compatibility
 - Easily combined with existing and future algorithms for RL
 - Better approximation of the state values
 - More frequent updating of the value stream
 - Only the value for one of the actions is updated in traditional deep Q network


Outline

- Introduction
- Deep Q Network
- Double DQN
- Dueling Network
- Experiment
- Conclusion

Experiment

- Compare with Marc et al. (2012)
- Achieving more than 75% of the human score on 29 games.

$$\text{SCORE}_{\text{normalized}} = \frac{\text{SCORE}_{\text{agent}} - \text{SCORE}_{\text{random}}}{\text{SCORE}_{\text{human}} - \text{SCORE}_{\text{random}}}$$




Experiment

- Replay memory and Target Q help to increase score.

Game	With replay, with target Q	With replay, without target Q	Without replay, with target Q	Without replay, without target Q
Breakout	316.8	240.7	10.2	3.2
Enduro	1006.3	831.4	141.9	29.1
River Raid	7446.6	4102.8	2867.7	1453.0
Seaquest	2894.4	822.6	1003.0	275.8
Space Invaders	1088.9	826.3	373.2	302.0


Experiment

- Double DQN vs DQN




Experiment

- Dueling Network vs Hasselt et al. (2015)



Experiment

- Dueling Network vs DDQN



Outline

- Introduction
- Deep Q Network
- Double DQN
- Dueling Network
- Experiment
- Conclusion

Conclusion

- DQN is the first successful Deep Reinforcement Learning Algorithm
 - Comparable level with human on 49 games of Atari2600
 - Publish on Nature 2015
- Double DQN solve the overestimation problem of DQN
 - Publish AAAI 2016
- Dueling Network is a new network architecture for RL
 - Compatibility for existed RL algorithm
 - Better performance
 - ICML 2016 Best Paper