

Deep Learning and Practice

Lab 10: Deep Deterministic Policy Gradient


李韓

0556157

fm.bigballon@gmail.com


Episode rewards

Reward



Smoothing 0.00

Reward



Smoothing 0.95

Explain the mechanism of critic updating

```
# loss & optimize op
self.loss = tflearn.mean_square(self.target_q_value, self.out)
self.optimize = tf.train.AdamOptimizer(self.learning_rate).minimize(self.loss)

def train(self, X, action, target_q_value):
    return self.sess.run([self.out, self.optimize], feed_dict={
        self.input: X,
        self.action: action,
        self.target_q_value: target_q_value
    })

# update critic
predicted_q_value, _ = behavior_critic.train(s_batch, a_batch, np.reshape(y_i, (MINIBATCH_SIZE, 1)))
```

Calculate the loss (MSE) $\left[\text{loss: } L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i | \theta^Q))^2 \right]$

Using AdamOptimizer to minimize the loss.

s_batch is the set of states,

a_batch is the set of action,

self.target_q_value is the set of y $\left[y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1} | \theta^{\mu'}) | \theta^{Q'}) \right]$

self.out is the predictions of critic network.

Describe the algorithm of actor updating

```
# -----
action = behavior_actor.predict(s_batch)
gradient = behavior_critic.action_gradients(s_batch, action)
behavior_actor.train(s_batch, gradient[0])

# soft update target networks
sess.run([actor_update_ops, critic_update_ops])
```

1. Use behavior_actor to predict actions $\left[\text{param: s_batch} \quad \text{return: action} \right]$

```
def predict(self, X):
    return self.sess.run(self.scaled_out, feed_dict={
        self.input: X
    })
```

2. Use behavior critic to compute action gradients [param: s_batch, action return:gradient]

```
def action_gradients(self, X, action):  
    return self.sess.run(self.action_grads, feed_dict={self.input: X, self.action: action})
```

compute the partial derivatives of self.out with respect to self.action

```
self.action_grads = tf.gradients(self.out, self.action)
```

3. Update behavior actor

```
def train(self, X, a_gradient):  
    self.sess.run(self.optimize, feed_dict={  
        self.input: X,  
        self.action_gradient: a_gradient  
    })
```

$$\nabla_{\theta^{\mu}} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a | \theta^Q) |_{s=s_i, a=\mu(s_i)} \nabla_{\theta^{\mu}} \mu(s | \theta^{\mu}) |_{s_i}$$

Using tf.gradients to combine the gradients.

```
self.actor_gradients = tf.gradients(self.scaled_out, self.network_params, -self.action_gradient)
```

Using apply_gradients to use the gradients


```
self.optimize = tf.train.AdamOptimizer(self.learning_rate).\  
    apply_gradients(zip(self.actor_gradients, self.network_params))
```

Update the target networks:

$$\begin{aligned}\theta^{Q'} &\leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'} \\ \theta^{\mu'} &\leftarrow \tau \theta^{\mu} + (1 - \tau) \theta^{\mu'}\end{aligned}$$

```
def soft_update_ops(sess, target_net, behavior_net):  
    update_ops = []  
    for behavior_v, target_v in zip(behavior_net.get_params(), target_net.get_params()):  
        # soft update  
        op = target_v.assign(TAU*behavior_v + (1.0 - TAU)*target_v)  
        update_ops.append(op)  
  
    return update_ops
```

Performance – Highest episode reward



One video during the training process

See attachments[video-1495603800.mp4].

The last one tensorflow checkpoint file

See attachments[checkpoint.zip].