

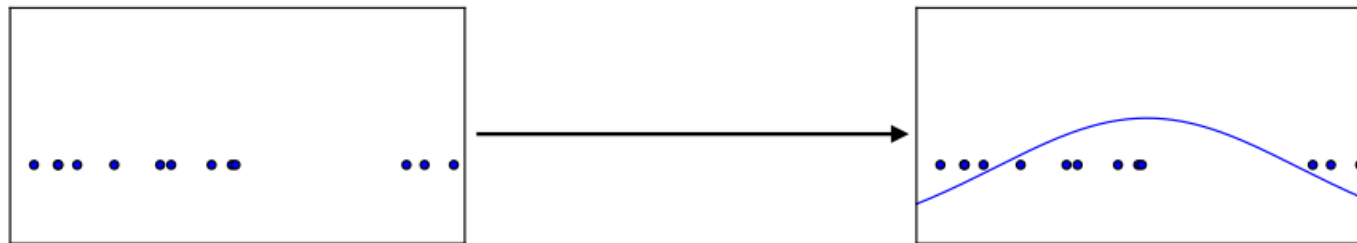
# Chapter 20

## Deep Generative Models

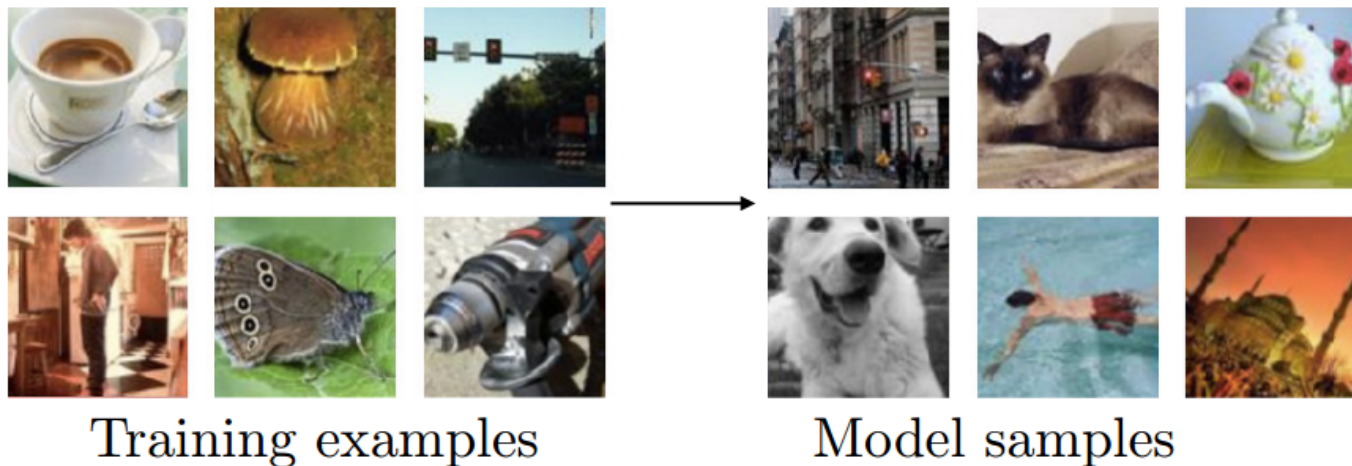
## Generative Models

Models that are able to

- Provide an estimate of the probability distribution function,  $p_{\text{data}}$ , or



- Generate samples from a (likely implicit) distribution

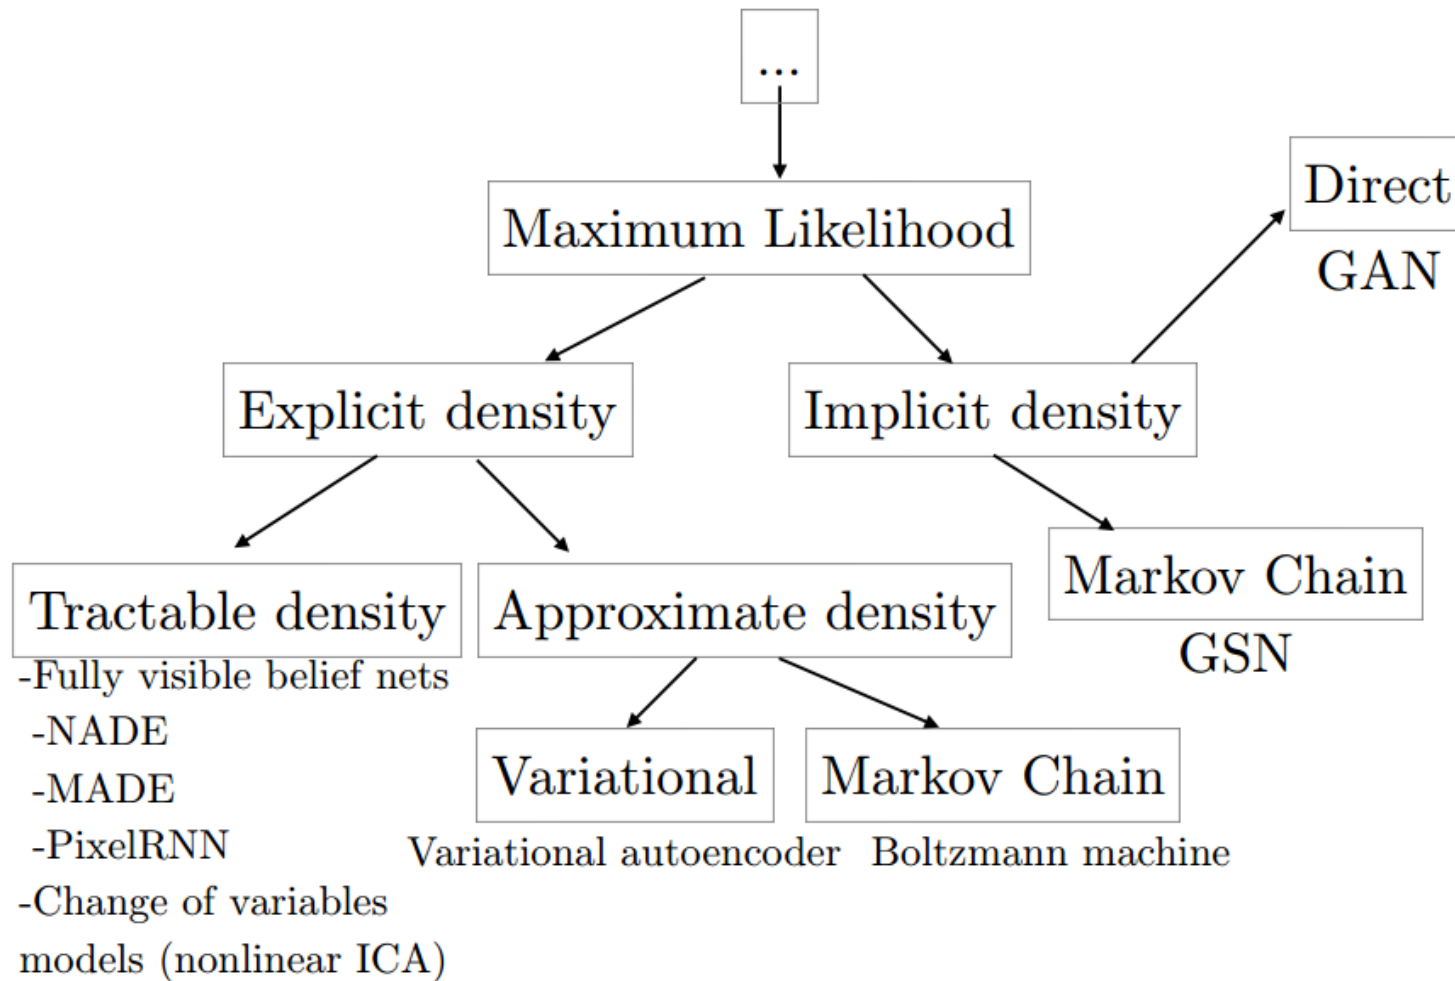


## Why Study Generative Models?

---

- Manipulation of high-dimensional, multi-modal distributions
- Potential uses in reinforcement learning, such as future state prediction
- Training with missing data (e.g. missing labels) and prediction on them
- Generation of realistic samples
- etc.


# Taxonomy of Generative Models



- Explicit density,  $p_{\text{model}}(\mathbf{x}; \boldsymbol{\theta})$ 
  - Tractable (trained with the ordinary ML)
  - Intractable/approximate (trained with approximate inference and/or MCMC approximations)
- Implicit density
  - Single-step sample generation via a network
  - Multi-step sample generation via Markov chains

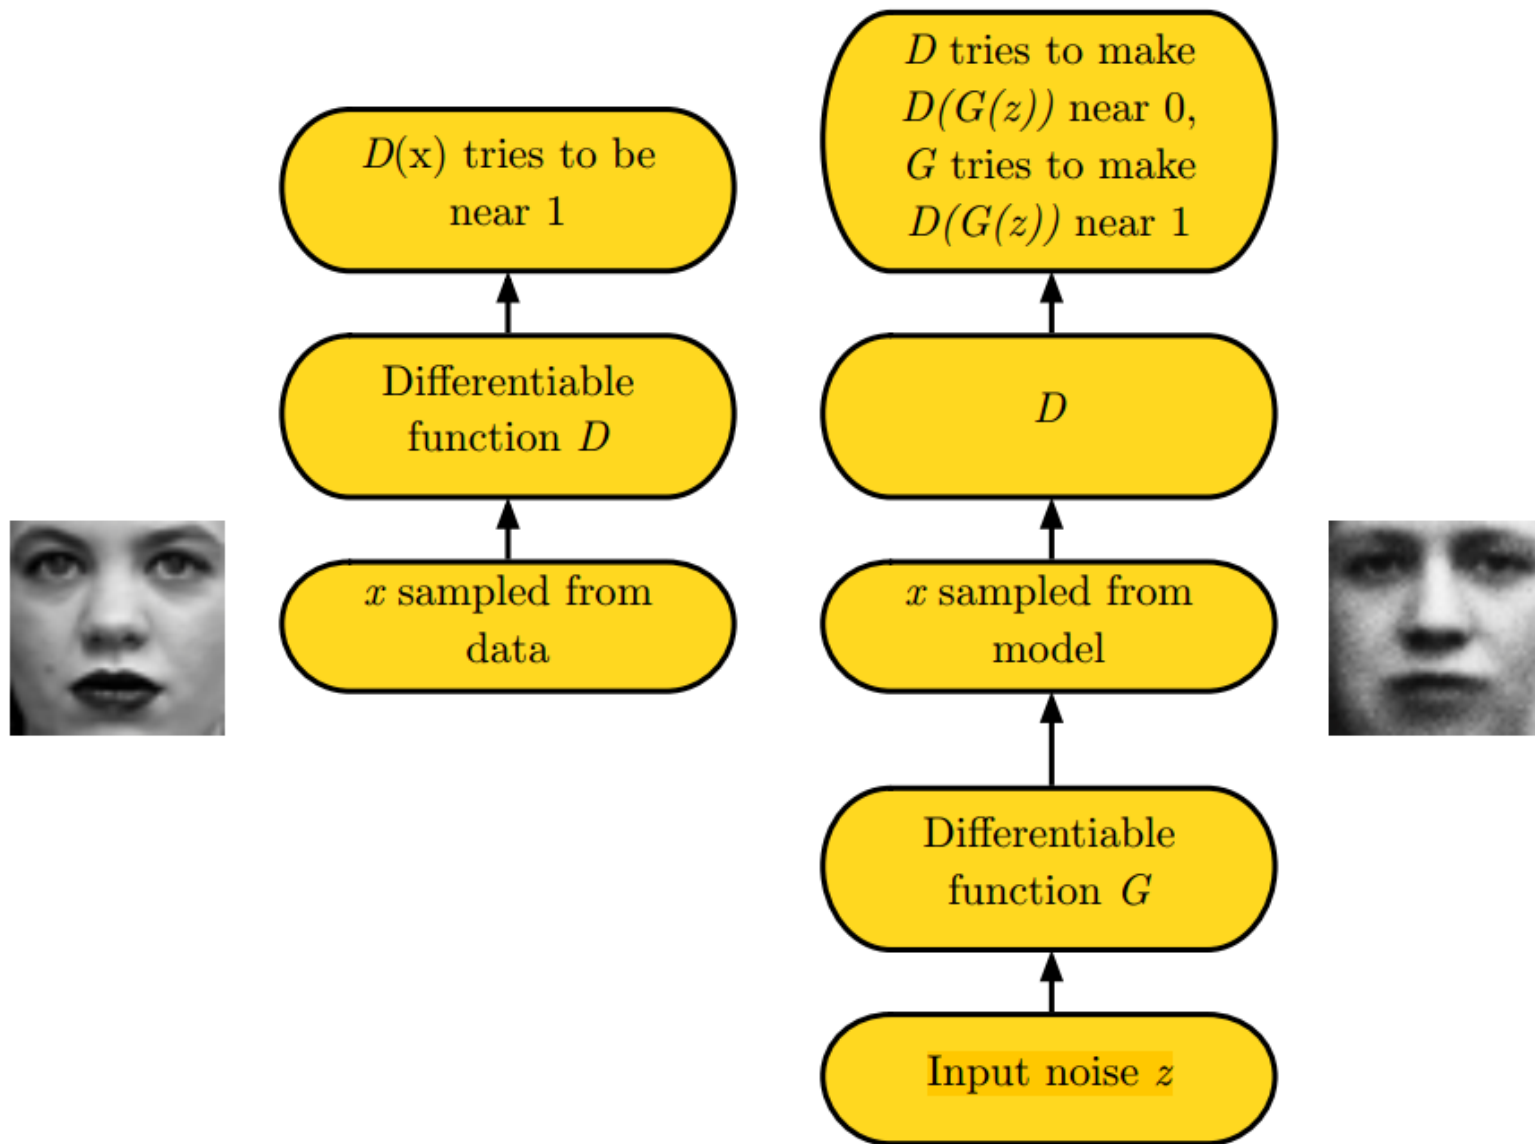
## Generative Adversarial Networks (GAN)

- A differentiable **generation** network  $G$ , paired with a **discriminator**  $D$  for training
- Generator  $G$  maps latent noises  $z \sim p(z)$  to visible variables  $x$ 
  - Conceptually, a graphical model with the same structure as VAE
  - $x = G(z)$  can be regarded as a sample drawn from some  $p_g(x)$



- **Generator is what we are concerned with**

- Discriminator  $D$  divides inputs into real and fake classes
  - An ordinary binary classifier trained supervisedly
  - Inputs are training examples (real) and generated samples (fake)





## Training GANs: Two-Player Minimax Game

---

- $D(\mathbf{x}; \boldsymbol{\theta}^{(D)})$ ,  $G(\mathbf{z}; \boldsymbol{\theta}^{(G)})$  are implemented with neural networks, and each has their own cost to minimize

- **Discriminator cost** (cross-entropy cost)

$$J^{(D)}(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)}) = -E_{\mathbf{x} \sim p_{\text{data}}} \log D(\mathbf{x}) - E_{\mathbf{z} \sim p_{\mathbf{z}}} \log(1 - D(G(\mathbf{z})))$$

where  $D(\mathbf{x})$  denotes the probability of  $\mathbf{x}$  being real

- **Generator cost**

$$J^{(G)}(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)}) = -J^{(D)}(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)})$$

- Note that the sum of all players' costs is zero (zero-sum game)

- The entire game can be summarized with a value function

$$V(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)}) \equiv -J^{(D)}(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)})$$

and the objective is to find a generator

$$\boldsymbol{\theta}^{(G)*} = \arg \min_{\boldsymbol{\theta}^{(G)}} \max_{\boldsymbol{\theta}^{(D)}} V(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)})$$

## Optimization vs. Game

- The solution to an optimization problem is generally a **local minimum** of an objective function in parameter space, e.g.

$$\arg \min_{\boldsymbol{\theta}^{(G)}, \boldsymbol{\theta}^{(D)}} V(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)})$$

where both  $\boldsymbol{\theta}^{(G)}, \boldsymbol{\theta}^{(D)}$  are optimized simultaneously

- The solution to a game problem is generally a **saddle point** of an objective function in parameter space, e.g.

$$\arg \min_{\boldsymbol{\theta}^{(G)}} \max_{\boldsymbol{\theta}^{(D)}} V(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)})$$

where  $\boldsymbol{\theta}^{(G)}, \boldsymbol{\theta}^{(D)}$  are optimized in turn by controlling one of them at a time with the other fixed

## The Optimal Discriminator

---

- For a given generator  $G$ , the optimal discriminator is seen to be

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})}$$

which can be obtained by having

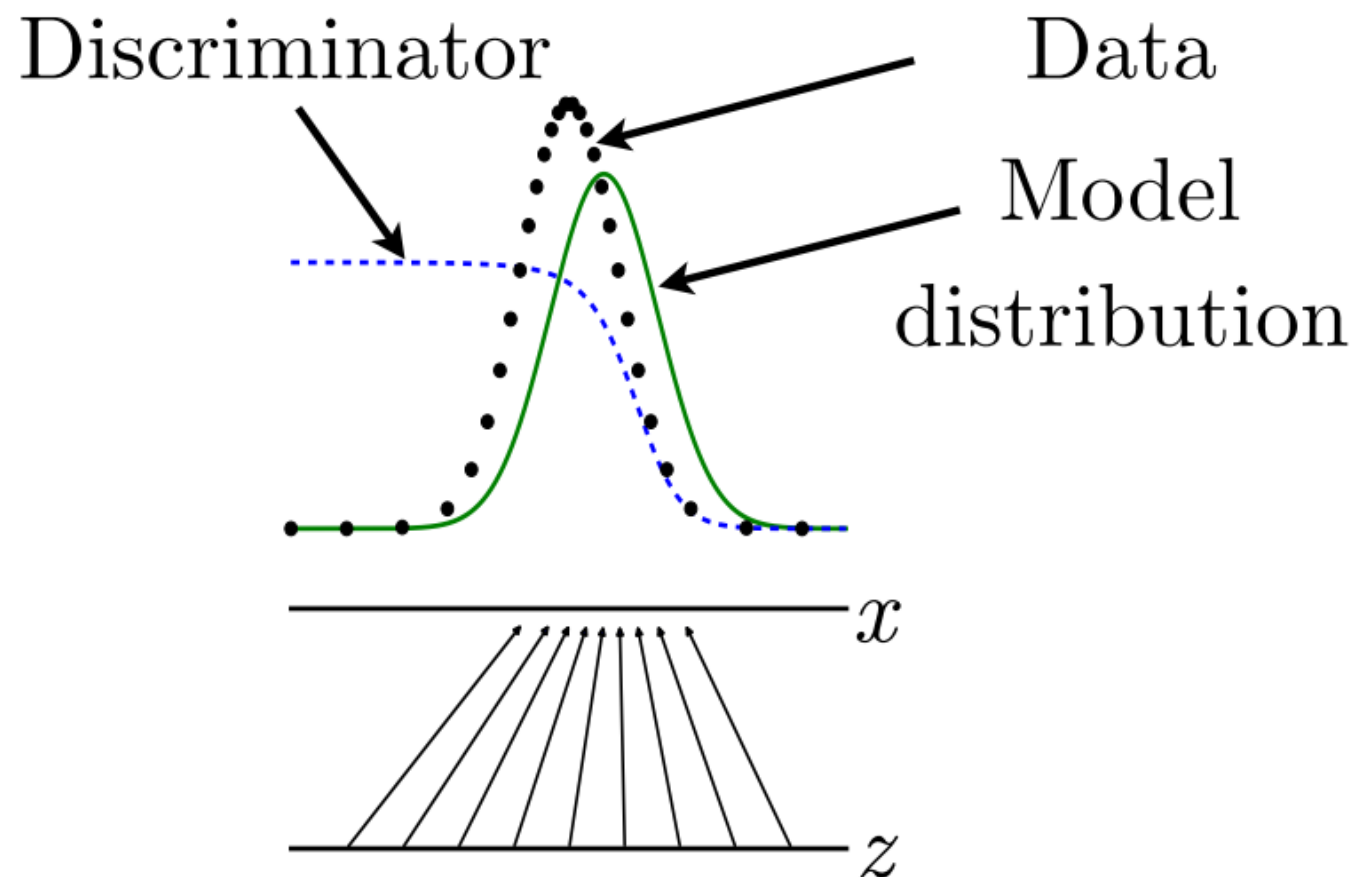
$$\frac{\delta}{\delta D(\mathbf{x})} J^{(D)}(\mathbf{x}) = 0$$

- When given enough capacity, the discriminator obtains an estimate

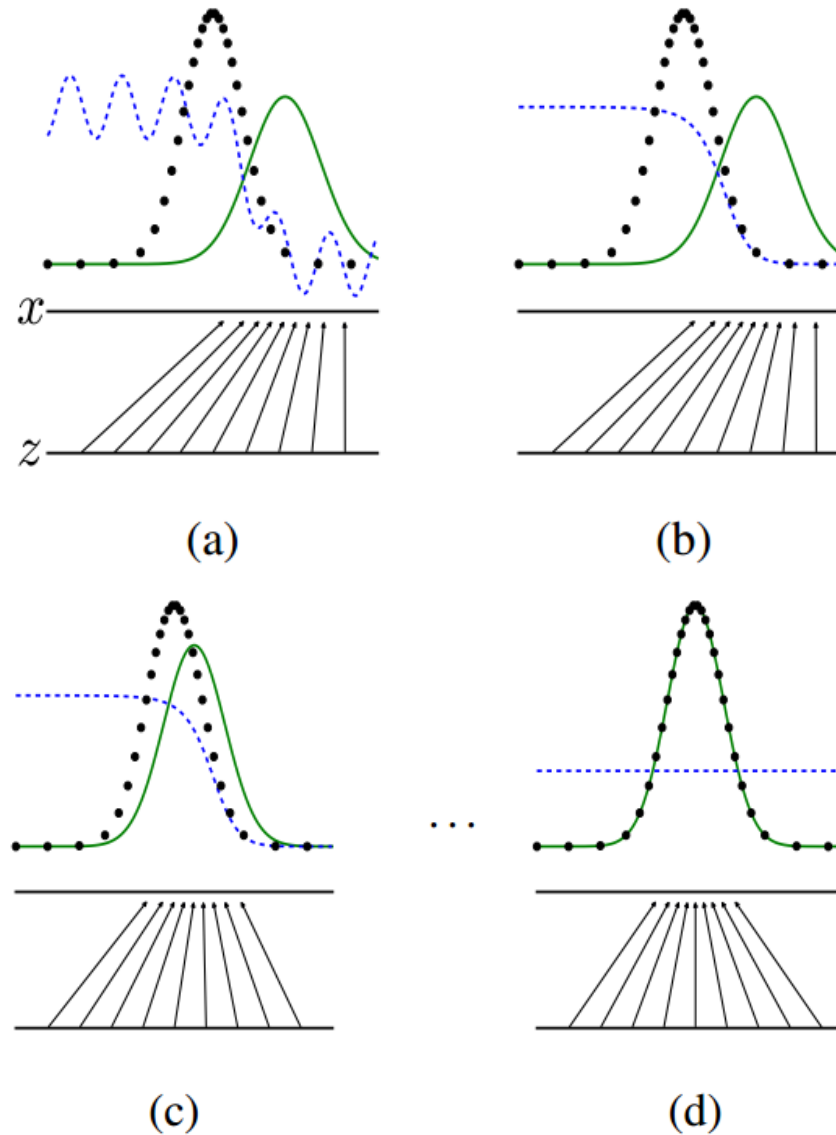
$$\frac{p_{\text{data}}(\mathbf{x})}{p_g(\mathbf{x})}$$

at every  $\mathbf{x}$

- This is the key that sets GANs apart from other generative models



- The generator is to learn a model by following a discriminator uphill



## The Optimal Generator

- Given  $D_G^*(\mathbf{x})$  and enough capacity, the optimal generator is to minimize the Jensen-Shannon divergence between  $p_{\text{data}}$  and  $p_g$

$$\begin{aligned}
 & \arg \min_{p_g} E_{\mathbf{x} \sim p_{\text{data}}} \log D_G^*(\mathbf{x}) + E_{\mathbf{x} \sim p_g} \log(1 - D_G^*(G(\mathbf{x}))) \\
 &= \arg \min_{p_g} E_{\mathbf{x} \sim p_{\text{data}}} \log \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})} + E_{\mathbf{x} \sim p_g} \log \frac{p_g(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})} \\
 &= \arg \min_{p_g} -\log(4) + \text{KL} \left( p_{\text{data}} \parallel \frac{p_{\text{data}} + p_g}{2} \right) + \text{KL} \left( p_g \parallel \frac{p_{\text{data}} + p_g}{2} \right) \\
 &= \arg \min_{p_g} -\log(4) + 2 \times \text{JSD}(p_{\text{data}} \parallel p_g)
 \end{aligned}$$

- The minimum is achieved when  $p_g = p_{\text{data}}$ , i.e.  $\text{JSD}(p_{\text{data}} \parallel p_g) = 0$

- Remarks

- The optimization is done w.r.t.  $p_g$  directly
- The analysis for the discriminator is done w.r.t.  $D(\mathbf{x})$
- Enough capacity in both contexts means that  $D_G^*(\mathbf{x})$  and  $p_g^*(\mathbf{x})$  can be implemented by  $D(\mathbf{x}; \boldsymbol{\theta}^{(D)*})$  and  $G(\mathbf{z}; \boldsymbol{\theta}^{(G)*})$ , respectively



## Implementation

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log \left( 1 - D(G(z^{(i)})) \right) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left( 1 - D(G(z^{(i)})) \right).$$

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

- (Convergence) If  $G$  and  $D$  have enough capacity, and at each step of Algorithm 1, the discriminator is allowed to reach its optimum  $D_G^*(\mathbf{x})$  given  $G$ , and  $p_g$  is updated to improve the criterion (reduce the cost)

$$E_{\mathbf{x} \sim p_{\text{data}}} \log D_G^*(\mathbf{x}) + E_{\mathbf{x} \sim p_g} \log(1 - D_G^*(G(\mathbf{x})))$$

then  $p_g$  converges to  $p_{\text{data}}$

- Nothing is said about the convergence when optimization is done based on simultaneous stochastic gradient descent in parameter space

## Non-Convergence of Gradient Descent

---

- Toy problem

$$\min_x \max_y V(x, y) = xy$$

- $x, y$  are optimized based on gradient descent with a tiny learning rate

$$x(t + \Delta t) = x(t) - \Delta t \frac{\partial}{\partial x(t)} V(x(t), y(t))$$

$$y(t + \Delta t) = y(t) + \Delta t \frac{\partial}{\partial y(t)} V(x(t), y(t))$$



- This amounts to solving

$$\begin{cases} x'(t) = -y(t) \\ y'(t) = x(t) \end{cases} \rightarrow x''(t) = -x(t)$$

which has a solution of the form

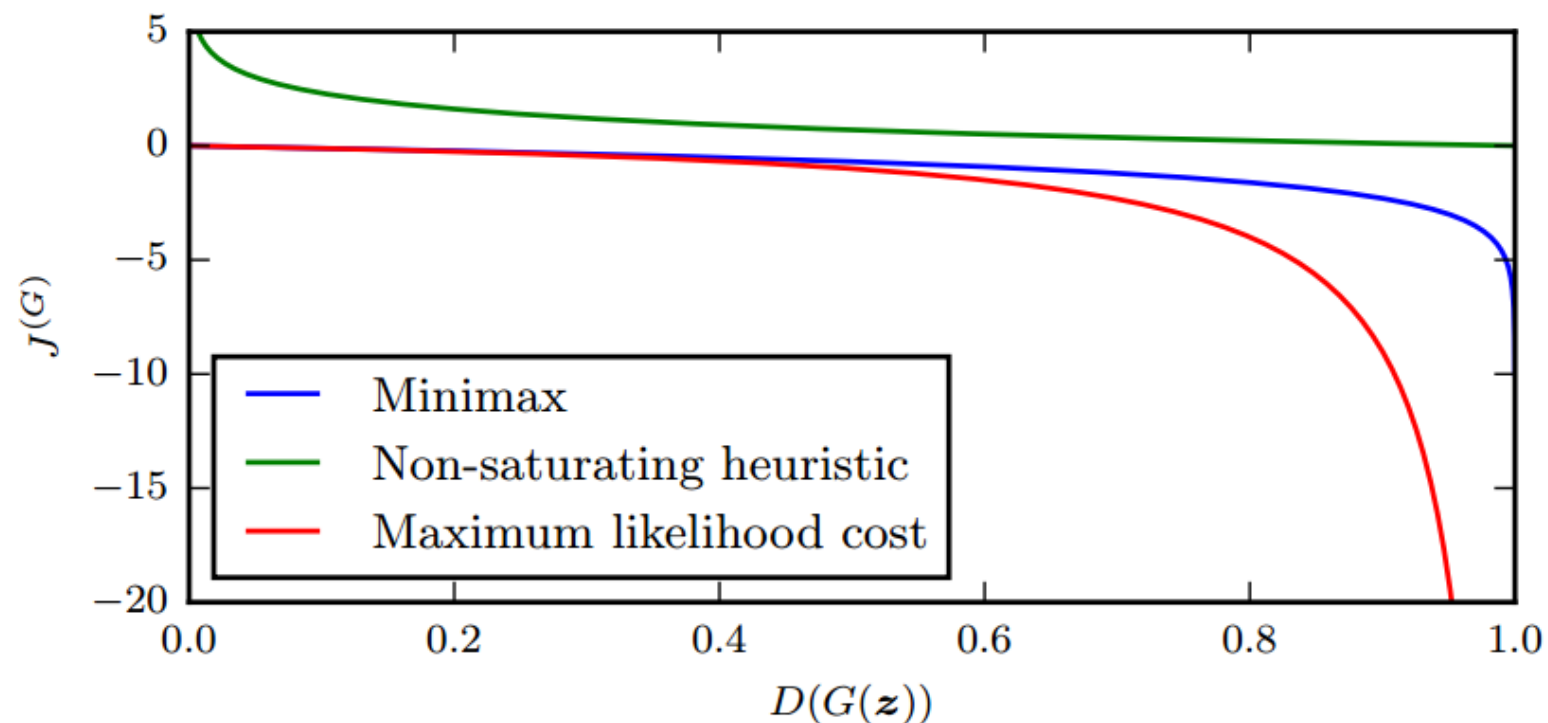
$$x(t) = x(0) \cos(t) + y(0) \sin(t)$$

$$y(t) = x(0) \sin(t) + y(0) \cos(t)$$



## Other Games

- Zero-sum game does not perform well in learning generator: gradients of  $J^{(G)}$  w.r.t  $D(G(z))$  vanish when the discriminator performs well



- Heuristic, non-saturating game (to ensure non-zero gradients)

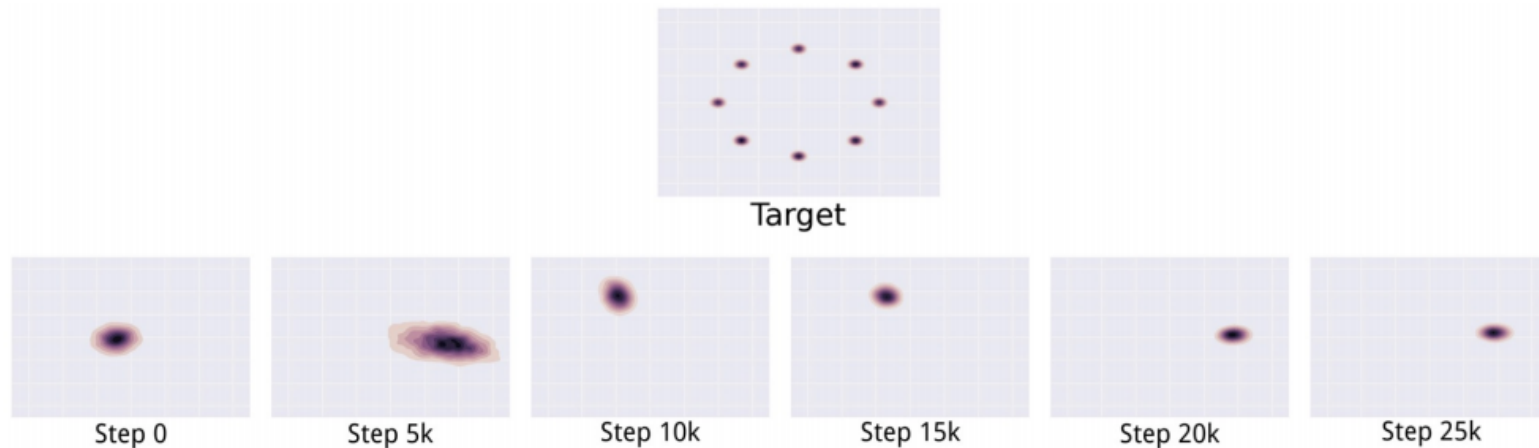
$$J^{(G)} = -E_z \log D(G(z))$$

- Maximum likelihood game (to minimize KL divergence)

$$J^{(G)} = -E_z \exp(\sigma^{-1}(D(G(z))))$$

## Mode Collapse Problem

- The generator learns to map different  $z$  to the same  $x$




Top: Data distribution (Mixture Gaussian)

Bottom: Learned generator distribution over time

- The generator distribution produces only a single mode at a time and does not converge in this example
- This is acceptable in some applications but not all

## Learned Representation


- The generator can learn a distributed representation that disentangles high-level concepts, e.g. gender vs. wearing glasses





# DCGAN

- There are many different implementations for generators, such as DCGAN, LPGAN, and more (study by yourself)



## Deep Boltzmann Machines (DBM)

---

- An energy-based generative model with an explicit density over **binary** visible  $\mathbf{v}$  and hidden  $\mathbf{h}^{(1)}, \mathbf{h}^{(2)}, \mathbf{h}^{(3)}$  variables

$$p(\mathbf{v}, \mathbf{h}^{(1)}, \mathbf{h}^{(2)}, \mathbf{h}^{(3)}; \boldsymbol{\theta}) = \frac{1}{Z(\boldsymbol{\theta})} \exp(-E(\mathbf{v}, \mathbf{h}^{(1)}, \mathbf{h}^{(2)}, \mathbf{h}^{(3)}; \boldsymbol{\theta}))$$

where


$$\begin{aligned} E(\mathbf{v}, \mathbf{h}^{(1)}, \mathbf{h}^{(2)}, \mathbf{h}^{(3)}; \boldsymbol{\theta}) \\ = -\mathbf{v}^T \mathbf{W}^{(1)} \mathbf{h}^{(1)} - \mathbf{h}^{(1)T} \mathbf{W}^{(2)} \mathbf{h}^{(2)} - \mathbf{h}^{(2)T} \mathbf{W}^{(3)} \mathbf{h}^{(3)} \end{aligned}$$

and

$$\boldsymbol{\theta} = \{\mathbf{W}^{(1)}, \mathbf{W}^{(2)}, \mathbf{W}^{(3)}\}$$

- Note that bias terms are omitted for simplicity

- Graphical model for DBM, where odd layers can be separated from even layers to reveal a bipartite structure



- As a result, variables in odd layers are conditionally independent given even layers and vice versa; this enables block Gibbs sampling

- Likewise, it is seen that variables in a layer are conditionally independent given the neighbouring layers
- In the case of two hidden layers, we have

$$p(v_i = 1 | \mathbf{h}^{(1)}) = \sigma(\mathbf{W}_{i,:}^{(1)} \mathbf{h}^{(1)})$$

$$p(h_i^{(1)} = 1 | \mathbf{v}, \mathbf{h}^{(2)}) = \sigma(\mathbf{v}^T \mathbf{W}_{:,i}^{(1)} + \mathbf{W}_{i,:}^{(2)} \mathbf{h}^{(2)})$$

$$p(h_i^{(2)} = 1 | \mathbf{h}^{(1)}) = \sigma(\mathbf{h}^{(1)T} \mathbf{W}_{:,i}^{(2)})$$

- However, the posterior distribution of all hidden layers given the visible layer does not factorize because of interactions between layers

$$p(\mathbf{h}^{(1)}, \mathbf{h}^{(2)} | \mathbf{v}) \neq \prod_j p(h_j^{(1)} | \mathbf{v}) \prod_k p(h_k^{(2)} | \mathbf{v})$$

- Approximate inference needs to be sought

## DBM Mean Field Inference

- To construct a factorial  $Q(\mathbf{h}|\mathbf{v})$  for approximating  $p(\mathbf{h}|\mathbf{v})$

$$p(\mathbf{h}^{(1)}, \mathbf{h}^{(2)}|\mathbf{v}) \approx Q(\mathbf{h}|\mathbf{v}) = \prod_j q(h_j^{(1)}|\mathbf{v}) \prod_k q(h_k^{(2)}|\mathbf{v})$$

- In the present case, all hidden variables  $h_j^{(1)}, h_k^{(2)}$  are binary; these  $q(h|\mathbf{v})$  must have a functional form of the Bernoulli distribution, i.e.

$$q(h_j^{(1)}|\mathbf{v}) = (\hat{h}_j^{(1)})^{h_j^{(1)}} (1 - \hat{h}_j^{(1)})^{(1-h_j^{(1)})}, \forall j$$

$$q(h_k^{(2)}|\mathbf{v}) = (\hat{h}_k^{(2)})^{h_k^{(2)}} (1 - \hat{h}_k^{(2)})^{(1-h_k^{(2)})}, \forall k$$

where  $\hat{h}_j^{(1)}, \hat{h}_k^{(2)} \in [0, 1]$  are the corresponding parameters

- Carrying out the expectation (needs some work)

$$\tilde{q}_j(h_j|\mathbf{v}) = \exp(E_{q_{-j}}(\log p(\mathbf{v}, \mathbf{h}^{(1)}, \mathbf{h}^{(2)}; \boldsymbol{\theta})))$$

yields the following fixed-point update equations

$$\hat{h}_j^{(1)} = \sigma \left( \sum_i v_i W_{i,j}^{(1)} + \sum_k W_{j,k}^{(2)} \hat{h}_k^{(2)} \right), \forall j$$

$$\hat{h}_k^{(2)} = \sigma \left( \sum_j W_{j,k}^{(2)} \hat{h}_j^{(1)} \right), \forall k$$

## DBM Parameter Learning

- DBM learning has to confront both the intractable inference  $p(\mathbf{h}|\mathbf{v})$  and the intractable partition function  $Z(\boldsymbol{\theta})$
- Combined variational inference, learning, and MCMC is necessary
- The objective then becomes to find  $\mathbf{W}^{(1)}, \mathbf{W}^{(2)}$  that minimize

$$\mathcal{L}(Q, \boldsymbol{\theta}) = \sum_i \sum_j v_i W_{i,j}^{(1)} \hat{h}_j^{(1)} + \sum_j \sum_k \hat{h}_j^{(1)} W_{j,k}^{(2)} \hat{h}_k^{(2)} - \log Z(\boldsymbol{\theta}) + H(Q)$$

which can be done via gradient descent

$$\boldsymbol{\theta}' = \boldsymbol{\theta} - \varepsilon \nabla_{\boldsymbol{\theta}} \mathcal{L}(Q, \boldsymbol{\theta})$$

(study Algorithm 20.1)

- In general, layer-wise pre-training is needed to arrive at a good model

---

Set  $\epsilon$ , the step size, to a small positive number

Set  $k$ , the number of Gibbs steps, high enough to allow a Markov chain of  $p(\mathbf{v}, \mathbf{h}^{(1)}, \mathbf{h}^{(2)}; \boldsymbol{\theta} + \epsilon \Delta_{\boldsymbol{\theta}})$  to burn in, starting from samples from  $p(\mathbf{v}, \mathbf{h}^{(1)}, \mathbf{h}^{(2)}; \boldsymbol{\theta})$ . Initialize three matrices,  $\tilde{\mathbf{V}}$ ,  $\tilde{\mathbf{H}}^{(1)}$  and  $\tilde{\mathbf{H}}^{(2)}$  each with  $m$  rows set to random values (e.g., from Bernoulli distributions, possibly with marginals matched to the model's marginals).

**while** not converged (learning loop) **do**

Sample a minibatch of  $m$  examples from the training data and arrange them as the rows of a design matrix  $\mathbf{V}$ .

Initialize matrices  $\hat{\mathbf{H}}^{(1)}$  and  $\hat{\mathbf{H}}^{(2)}$ , possibly to the model's marginals.

**while** not converged (mean field inference loop) **do**

$$\hat{\mathbf{H}}^{(1)} \leftarrow \sigma \left( \mathbf{V} \mathbf{W}^{(1)} + \hat{\mathbf{H}}^{(2)} \mathbf{W}^{(2)\top} \right).$$

$$\hat{\mathbf{H}}^{(2)} \leftarrow \sigma \left( \hat{\mathbf{H}}^{(1)} \mathbf{W}^{(2)} \right).$$

**end while**

$$\Delta \mathbf{W}^{(1)} \leftarrow \frac{1}{m} \mathbf{V}^\top \hat{\mathbf{H}}^{(1)}$$

$$\Delta \mathbf{W}^{(2)} \leftarrow \frac{1}{m} \hat{\mathbf{H}}^{(1)\top} \hat{\mathbf{H}}^{(2)}$$

**for**  $l = 1$  to  $k$  (Gibbs sampling) **do**

Gibbs block 1:

$$\forall i, j, \tilde{V}_{i,j} \text{ sampled from } P(\tilde{V}_{i,j} = 1) = \sigma \left( \mathbf{W}_{j,:}^{(1)} \left( \tilde{\mathbf{H}}_{i,:}^{(1)} \right)^\top \right).$$

$$\forall i, j, \tilde{H}_{i,j}^{(2)} \text{ sampled from } P(\tilde{H}_{i,j}^{(2)} = 1) = \sigma \left( \tilde{\mathbf{H}}_{i,:}^{(1)} \mathbf{W}_{:,j}^{(2)} \right).$$

Gibbs block 2:

$$\forall i, j, \tilde{H}_{i,j}^{(1)} \text{ sampled from } P(\tilde{H}_{i,j}^{(1)} = 1) = \sigma \left( \tilde{\mathbf{V}}_{i,:} \mathbf{W}_{:,j}^{(1)} + \tilde{\mathbf{H}}_{i,:}^{(2)} \mathbf{W}_{j,:}^{(2)\top} \right).$$

**end for**

$$\Delta \mathbf{W}^{(1)} \leftarrow \Delta \mathbf{W}^{(1)} - \frac{1}{m} \mathbf{V}^\top \tilde{\mathbf{H}}^{(1)}$$

$$\Delta \mathbf{W}^{(2)} \leftarrow \Delta \mathbf{W}^{(2)} - \frac{1}{m} \tilde{\mathbf{H}}^{(1)\top} \tilde{\mathbf{H}}^{(2)}$$

$\mathbf{W}^{(1)} \leftarrow \mathbf{W}^{(1)} + \epsilon \Delta \mathbf{W}^{(1)}$  (this is a cartoon illustration, in practice use a more effective algorithm, such as momentum with a decaying learning rate)

$$\mathbf{W}^{(2)} \leftarrow \mathbf{W}^{(2)} + \epsilon \Delta \mathbf{W}^{(2)}$$

**end while**

---



## Topics Not Covered

---

- Optimization for training deep models (Chapter 8)
- Representation learning (Chapter 15)
- Back-prop through random operations (REINFORCE, Chapter 20)
- BM for real-valued data (Chapter 20)
- Generative Stochastic Networks (Chapter 20)
- Deep Belief Networks (Chapter 20)
- Other generative models (Chapter 20)