

Markov Decision Process

I-Chen Wu

- Sutton, R.S. and Barto, A.G., Reinforcement Learning: An Introduction, MIT Press, Cambridge, MA, 1998. (Bible for RL)
 - <http://webdocs.cs.ualberta.ca/~sutton/book/ebook/the-book.html>
 - Chapters 3-4
- David Silver, Online Course for Deep Reinforcement Learning.
 - <http://www.cs.ucl.ac.uk/staff/D.Silver/web/Teaching.html>
 - Chapters 2-3



Outline

- Introduction
- Markov Property
- Markov Process
- Markov Reward Process (MRP)
- Markov Decision Process (MDP)
- Partially Observable Markov Decision Process (POMDP)

The purpose of this chapter:

- Introduce all kinds of Markov processes

Introduction

- Markov decision processes formally describe an environment for reinforcement learning
 - where the environment is fully observable.
 - i.e. The current state completely characterizes the process
 - E.g., 2048.
- Almost all RL problems can be formalized as MDPs, e.g.
 - Optimal control primarily deals with continuous MDPs
 - Partially observable problems can be converted into MDPs
 - Bandits are MDPs with one state



Markov Property

- Markov Property:

- “The future is independent of the past given the present”
- Definition: A state S_t is Markov if and only if
$$\mathbb{P}[S_{t+1} | S_t] = \mathbb{P}[S_{t+1} | S_1, \dots, S_t]$$

- Comments:

- The state captures all relevant information from the history
- Once the state is known, the history may be thrown away
- i.e. The state is a sufficient statistic of the future

- But, what if the history does matter?

- Simply let S_t carry all information of history, $H_t = (S_1, \dots, S_{t-1})$.
 - ▶ E.g., the castling rule for chess.
- Then, it satisfies Markov Property.



Markov Process

- A Markov process is a memoryless random process,
 - i.e. a sequence of random states S_1, S_2, \dots with the Markov property.

Definition:

- A Markov Process (or Markov Chain) is a tuple $\langle \mathcal{S}, \mathcal{P} \rangle$
 - \mathcal{S} is a (finite) set of states
 - \mathcal{P} is a state transition probability matrix (part of the environment),
$$\mathcal{P}_{ss'} = \mathbb{P}[S_{t+1} = s' \mid S_t = s]$$

State Transition Matrix

- For a Markov state s and successor state s' , the **state transition probability** is defined by

$$\mathcal{P}_{ss'} = \mathbb{P}[S_{t+1} = s' \mid S_t = s]$$

- State transition matrix \mathcal{P} : (assume n states)

$$\mathcal{P} = \begin{bmatrix} \mathcal{P}_{11} & \dots & \mathcal{P}_{1n} \\ \dots & & \dots \\ \mathcal{P}_{n1} & \dots & \mathcal{P}_{nn} \end{bmatrix}$$

- Each row of matrix sums to 1.

Markov Reward Process (MRP)


- A Markov reward process is a Markov chain with values.

Definition:

- A **Markov Reward Process** is a tuple $\langle \mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
 - \mathcal{S} is a finite set of states
 - \mathcal{P} is a state transition probability matrix (part of the environment),
 $\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$
 - \mathcal{R} is a reward function,
 $\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$
 - γ is a discount factor $\gamma \in [0, 1]$.



Example: Student MRP



Return

Definition

- The return G_t is the total discounted reward from time-step t .

$$G_t = R_{t+1} + \gamma R_{t+2} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Notes:

- The discount $\gamma \in [0, 1]$ is the present value of future rewards
- The value of receiving reward R is diminishing
 - $\gamma^k R$, after $k + 1$ time-steps.
- This values immediate reward above delayed reward.
- Discount:
 - γ close to 0 leads to "myopic" evaluation
 - γ close to 1 leads to "far-sighted" evaluation



Value Function

- The value function $v(s)$ gives the long-term value of s
- Definition
 - The state value function $v(s)$ of an MRP is the expected return starting from state s
 - $v(s) = \mathbb{E}[G_t | S_t = s]$

Bellman Equation for MRPs

- The value function can be decomposed into two parts:
 - immediate reward R_{t+1}
 - discounted value of successor state $\gamma v(S_{t+1})$
- $$\begin{aligned} v(s) &= \mathbb{E}[G_t \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \cdots) \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) \mid S_t = s] \end{aligned}$$
- For a transition (s, r, s') , we have

$$v(s) = \mathcal{R}_s + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'} v(s')$$



Bellman Equation in Matrix Form

- The Bellman equation can be expressed concisely using matrices,

$$v = \mathcal{R} + \gamma \mathcal{P}v$$

- where v is a column vector with one entry per state.

$$\begin{bmatrix} v(1) \\ \dots \\ v(n) \end{bmatrix} = \begin{bmatrix} R_1 \\ \dots \\ R_n \end{bmatrix} + \gamma \begin{bmatrix} \mathcal{P}_{11} & \dots & \mathcal{P}_{1n} \\ \dots & \dots & \dots \\ \mathcal{P}_{n1} & \dots & \mathcal{P}_{nn} \end{bmatrix} \begin{bmatrix} v(1) \\ \dots \\ v(n) \end{bmatrix}$$

Solving the Bellman Equation

- The Bellman equation is a linear equation
- It can be solved directly:
$$v = \mathcal{R} + \gamma \mathcal{P}v$$
$$v = (1 - \gamma \mathcal{P})^{-1} \mathcal{R}$$
- Computational complexity is $O(n^3)$ for n states
- Direct solution only possible for small MRPs
- There are many iterative methods for large MRPs, e.g.
 - Dynamic programming
 - Monte-Carlo evaluation
 - Temporal-Difference learning

Markov Decision Processes (MDP)

- A Markov Decision Process is a tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
 - \mathcal{S} is a finite set of states
 - \mathcal{A} is a finite set of actions
 - \mathcal{P} is a state transition probability matrix (part of the environment),
$$\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$$
 - \mathcal{R} is a reward function,
$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$$
 - γ is a discount factor $\gamma \in [0, 1]$.

Example: Recycling Robot




Figure 3.3: Transition graph for the recycling robot example.



Example: Recycling Robot

● Transition and Rewards:

| s | s' | a | $p(s' s, a)$ | $r(s, a, s')$ |
|------|------|----------|--------------|---------------------|
| high | high | search | α | r_{search} |
| high | low | search | $1 - \alpha$ | r_{search} |
| low | high | search | $1 - \beta$ | -3 |
| low | low | search | β | r_{search} |
| high | high | wait | 1 | r_{wait} |
| high | low | wait | 0 | r_{wait} |
| low | high | wait | 0 | r_{wait} |
| low | low | wait | 1 | r_{wait} |
| low | high | recharge | 1 | 0 |
| low | low | recharge | 0 | 0. |



Policies

- A policy is the agent's behavior
 - It is a map from state to action
 - A policy fully defines the behaviour of an agent
 - MDP policies depend on the current state (not the history)
 - ▶ i.e. Policies are stationary (time-independent),
 $A_t \sim \pi(\cdot | S_t), \forall t > 0$
- Policy types:
 - **Deterministic policy:** $a = \pi(s_i)$
 - **Stochastic policy:** $\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$
 - ▶ Sometimes, written in $\pi(s, a)$.
- Examples:
 - In 2048: Up/down/left/right
 - In robotics: angle/force/...



Policy and MRP

- Given an MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ and a policy π
- The state sequence S_1, S_2, \dots is a Markov process $\langle \mathcal{S}, \mathcal{P}^\pi \rangle$
- The state and reward sequence $S_1, R_2, S_2, R_3, \dots$ becomes a Markov reward process (MRP) $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$
 - \mathcal{P}^π is a state transition probability matrix (part of the environment),

$$\mathcal{P}_{ss'}^\pi = \sum_{a \in \mathcal{A}} \pi(a|s) \mathcal{P}_{ss'}^a$$

- \mathcal{R}^π is a reward function,

$$\mathcal{R}_s^\pi = \sum_{a \in \mathcal{A}} \pi(a|s) \mathcal{R}_s^a$$

- So, the property of MRP can be applied.


Value Function

- A value function is a prediction of future reward
 - Used to evaluate the goodness/badness of states
 - ▶ therefore to select between actions.
 - Return $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$
- Types of value functions under policy π :
 - State value function: the expected return from s .
$$\begin{aligned} v_{\pi}(s) &= \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s] \\ &= \mathbb{E}_{\pi}[G_t \mid S_t = s] \end{aligned}$$
 - Q-Value function: the expected return from s taking action a .
$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a]$$
- Examples:
 - In 2048, the expected score from a board S_t .

Bellman Expectation Equation for π

- State value function:

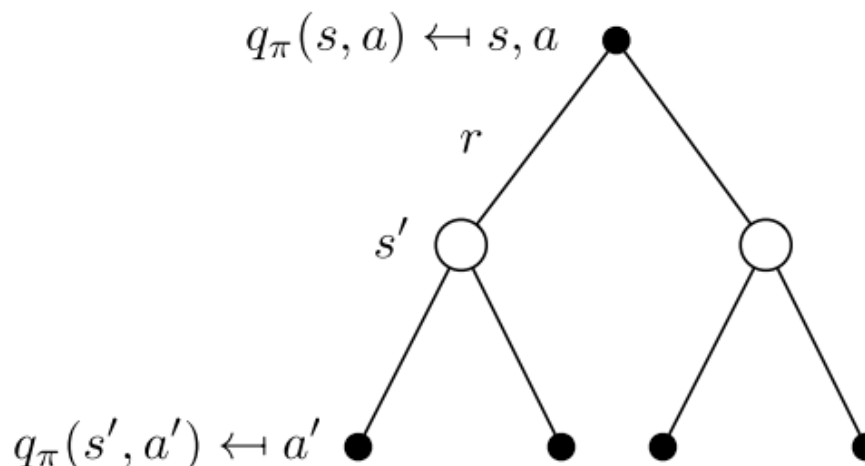
$$v_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left(\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_{\pi}(s') \right)$$



Bellman Expectation Equation for π

● Q value

$$q_{\pi}(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a'|s') q_{\pi}(s', a')$$



Bellman Expectation Equation in Matrix

- The Bellman expectation equation can be expressed concisely using the induced MRP.
- So, it can be solved directly:

$$v_{\pi} = \mathcal{R}^{\pi} + \gamma \mathcal{P}^{\pi} v_{\pi}$$
$$v_{\pi} = (1 - \gamma \mathcal{P}^{\pi})^{-1} \mathcal{R}^{\pi}$$

Optimal Value Function

- The optimal state-value function $v_*(s)$ is the maximum value function over all policies

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

- The optimal action-value function $q_*(s, a)$ is the maximum action-value function over all policies

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

- Notes:

- The optimal value function specifies the best possible performance in the MDP.
- An MDP is “solved” when we know the optimal value fn.



Optimal Policy

- Define a partial ordering over policies

$$\pi \geq \pi', v_*(s) = \max_{\pi} v_{\pi}(s) \text{ if } v_{\pi}(s) \geq v_{\pi'}(s), \forall s$$

- Theorem:


- There exists an optimal policy π_* that is better than or equal to all other policies, $\pi_* \geq \pi, \forall \pi$.
- All optimal policies achieve the optimal value function,
- All optimal policies achieve the optimal action-value function,

Finding an Optimal Policy

- An optimal policy can be found by maximizing over $q_*(s, a)$,
 - $\pi(a|s) = 1$, if $a = \operatorname{argmax}_a q_*(s, a)$
 - $\pi(a|s) = 0$, otherwise.
- There is always a deterministic optimal policy for any MDP
- If we know $q_*(s, a)$, we immediately have the optimal policy
- What about state value function $v_*(s)$?
 - Similar, but we need to know model, $\mathcal{P}_{ss'}^a$. \rightarrow not model free.




Bellman Optimality Equation for V^*



$$v_*(s) = \max_a \left(\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s') \right)$$



Bellman Optimality Equation for Q^*



$$q_\pi(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \max_{a' \in \mathcal{A}} q_\pi(s, a')$$



Solving the Bellman Optimality Equation

- Bellman Optimality Equation is non-linear
- No closed form solution (in general)
- Many iterative solution methods
 - Value Iteration
 - Policy Iteration
 - Q-learning
 - Sarsa
- Extensions to MDPs
 - Infinite and continuous MDPs
 - Partially observable MDPs
 - Undiscounted, average reward MDPs



Partially Observable Markov Decision Processes (POMDP)

- A POMDP is a tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{O}, \mathcal{P}, \mathcal{R}, \mathcal{Z}, \gamma \rangle$
 - \mathcal{S} is a finite set of states
 - \mathcal{A} is a finite set of actions
 - \mathcal{O} is a finite set of observations
 - \mathcal{P} is a state transition probability matrix,
 $\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$
 - \mathcal{R} is a reward function,
 $\mathcal{R}_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$
 - \mathcal{Z} is an observation function,
 $\mathcal{Z}_{s'o}^a = \mathbb{P}[O_{t+1} = o | S_{t+1} = s', A_t = a]$
 - γ is a discount factor $\gamma \in [0, 1]$.
- Examples:
 - Like Mahjong (麻將) in games, hidden scene in robotics.
- POMDP vs. MDP
 - A MDP can be trivially mapped onto a POMDP
 - A POMDP can be mapped onto an MDP:




Example: Mahjong

- Partially observable:



Value Iteration



$$v_{n+1}(s) = \max_{a \in A} \left(\mathcal{R}_s^a + \gamma \sum_{s' \in S} \mathcal{P}_{ss'}^a v_n(s') \right)$$

or:

$$V^{(n+1)}(s) = \max_{a \in \mathcal{A}} \left(\mathbb{E}_{s'|s,a} [r + \gamma V^{(n)}(s')] \right)$$



Operator View

- Value iteration update is viewed as:

- A function $\mathcal{T}: \mathcal{S} \rightarrow \mathcal{S}$.
- Called **backup operator**.

$$[\mathcal{T}V](s) = \max_{a \in \mathcal{A}} (\mathbb{E}_{s'|s,a} [r + \gamma V(s')])$$

Algorithm Value Iteration

Initialize $V(0)$ arbitrarily.

for $n = 0, 1, 2, \dots$ until termination condition do

$$V^{(n+1)} = \mathcal{T}V^{(n)}$$

end

Contraction Mapping View

- Backup operator \mathcal{T} is a contraction with modulus γ under ∞ -norm

$$\|\mathcal{T}V - \mathcal{T}W\|_{\infty} \leq \gamma \|V - W\|_{\infty}$$

- By contraction-mapping principle, it has a fixed point V^*
 - by iterating

$$V, \mathcal{T}V, \mathcal{T}^2V, \dots \rightarrow V^*$$



Policy Evaluation

- Problem: how to evaluate fixed policy π :

$$V^\pi(s) = \mathbb{E}_\pi[G_t | S_t = s] = \mathbb{E}_\pi[R_{t+1} + \gamma V^\pi(S_{t+1}) | S_t = s]$$

- Backwards recursion involves a backup operation

$$V^{(k+1)} = \mathcal{T}^\pi V^{(k)}$$

- \mathcal{T}^π is defined as:

$$[\mathcal{T}^\pi V](s) = \mathbb{E}_{s'|s, a=\pi(s)}[r + \gamma V(s')]$$

- \mathcal{T}^π is also a contraction with modulus γ , sequence

$$V, \mathcal{T}^\pi V, (\mathcal{T}^\pi)^2 V, (\mathcal{T}^\pi)^3 V, \dots \rightarrow V^\pi$$

- $V = \mathcal{T}^\pi V$ is a linear equation that we can solve directly.



Policy Iteration: Overview

- Alternate between
 - Evaluate policy $\pi \Rightarrow V^\pi$
 - Set new policy to be greedy policy for V^π
$$\pi(s) = \underset{a}{\operatorname{argmax}} \mathbb{E}_{s'|s,a} [R_{t+1} + \gamma V^\pi(s')]]$$
- Guaranteed to converge to optimal policy and value function in a finite number of iterations, when $\gamma < 1$
- Value function converges faster than in value iteration

Algorithm Policy Iteration

Initialize $\pi^{(0)}$ arbitrarily.

for $n = 1, 2, \dots$ until termination condition do

$$V^{(n+1)} = \text{Solve } [V = \mathcal{T}^{\pi^{(n-1)}} V]$$

$$\pi^{(n)} = \mathcal{G} V^{(n)} \quad \mathcal{G}: \text{a greedy mapping function.}$$

end



Modified Policy Iteration

- Update π to be the greedy policy, then value function with k backups (k -step lookahead)

Algorithm Modified Policy Iteration

Initialize $V^{(0)}$ arbitrarily.

for $n = 1, 2, \dots$ until termination condition do

$$\pi^{(n+1)} = \mathcal{G}V^{(n)}$$

$$V^{(n+1)} = (\mathcal{T}^{\pi^{n+1}})^k V^{(n)}, \text{ for integer } k \geq 1.$$

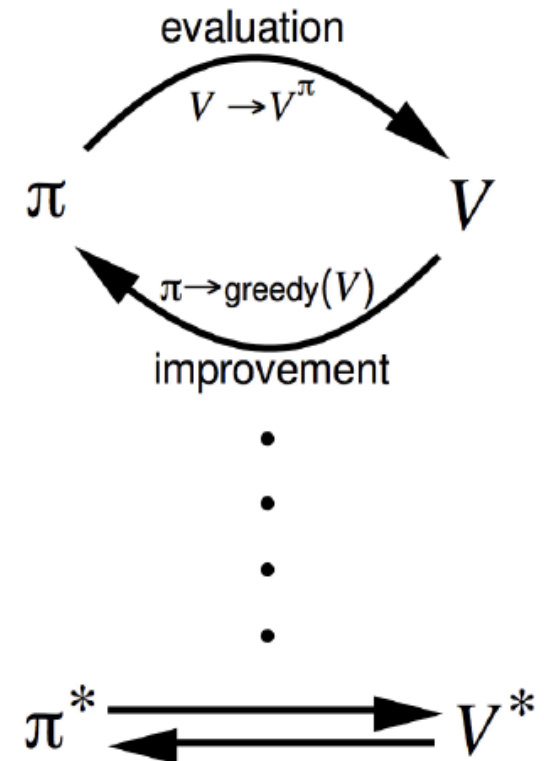
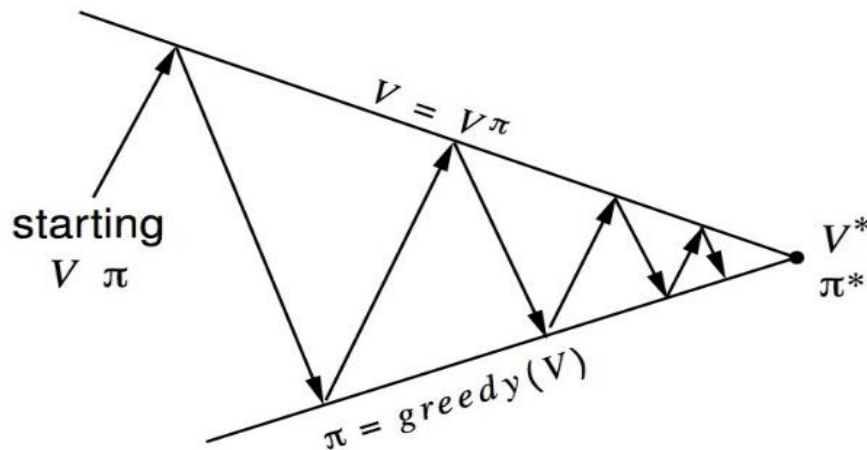
end

- $k = 1$: value iteration
- $k = \infty$: policy iteration



Policy Iteration in Different View

- Policy evaluation \rightarrow Estimate v_π
 - Iterative policy evaluation
- Policy improvement \rightarrow Generate $\pi' \geq \pi$
 - Greedy policy improvement



Policy Improvement

- Definition of **policy improvement**
 - Let π and π' be any pair of deterministic policies
 - ▶ For all $s \in S$, “ $\pi(s)$ performs better than $\pi'(s)$ ”
 - ▶ or $\pi \geq \pi'$.
- Given a policy π
 - Evaluate the policy π
$$v_\pi(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \dots | S_t = s]$$
 - **Improve the policy** by acting greedily with respect to v_π
$$\pi' = \text{greedy}(v_\pi)$$
- Notes:
 - In Small Gridworld improved policy was optimal, $\pi' = \pi^*$
 - In general, need more iterations of improvement / evaluation
 - But this process of policy iteration always converges to π^*



Proof of Policy Improvement

- Consider a deterministic policy, $a = \pi(s)$
- We can improve the policy by acting greedily
$$\pi'(s) = \operatorname{argmax}_{a \in A} q_{\pi}(s, a)$$
- This improves the value from any state s over one step,
$$q_{\pi}(s, \pi'(s)) = \max_{a \in A} q_{\pi}(s, a) \geq q_{\pi}(s, \pi(s)) = v_{\pi}(s)$$
- It therefore improves the value function, $v_{\pi'}(s) \geq v_{\pi}(s)$.

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)) = \mathbb{E}_{\pi'}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 q_{\pi}(S_{t+2}, \pi'(S_{t+2})) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \cdots | S_t = s] \\ &= v_{\pi'}(s) \end{aligned}$$



Convergence of Policy Improvement

- If improvements stop,

$$q_{\pi}(s, \pi'(s)) = \max_{a \in A} q_{\pi}(s, a) \geq q_{\pi}(s, \pi(s)) = v_{\pi}(s)$$

- Then the Bellman optimality equation has been satisfied

$$v_{\pi}(s) = \max_{a \in A} q_{\pi}(s, a)$$

- This implies $v_{\pi}(s) = v_*(s)$ for all $s \in S$

- The above proves that π will converge to an optimal policy.