

Model Free Reinforcement Learning

I-Chen Wu

- Sutton, R.S. and Barto, A.G., Reinforcement Learning: An Introduction, MIT Press, Cambridge, MA, 1998. (Bible for RL)
 - <http://webdocs.cs.ualberta.ca/~sutton/book/ebook/the-book.html>
 - Chapters 5-6
- David Silver, Online Course for Deep Reinforcement Learning.
 - <http://www.cs.ucl.ac.uk/staff/D.Silver/web/Teaching.html>
 - Chapters 4-5



Outline

- **Model-free prediction:** Estimate the value function of an unknown MDP
 - Monte-Carlo (MC) Learning
 - Temporal Difference (TD) Learning
- **Model-free control:** Approximate optimal policies based on the estimation of the value function.
 - On-Policy Monte-Carlo Control
 - On-Policy Temporal-Difference Learning
 - Off-Policy Learning

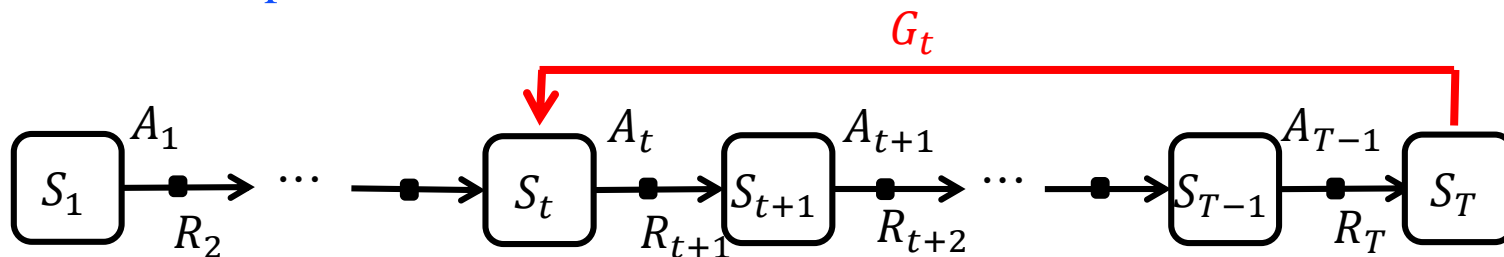
The purpose of this chapter:

- **Learn model-free RL: MC + TD**



Monte-Carlo Reinforcement Learning

- MC methods learn **directly from episodes of experience**
- MC is **model-free**:
 - no knowledge of MDP transitions / rewards
- MC learns from **complete episodes**:
 - no bootstrapping
- MC uses the simplest possible idea:
 - **value = mean return**
- Caveat: can only apply MC to episodic MDPs
 - **All episodes must terminate**

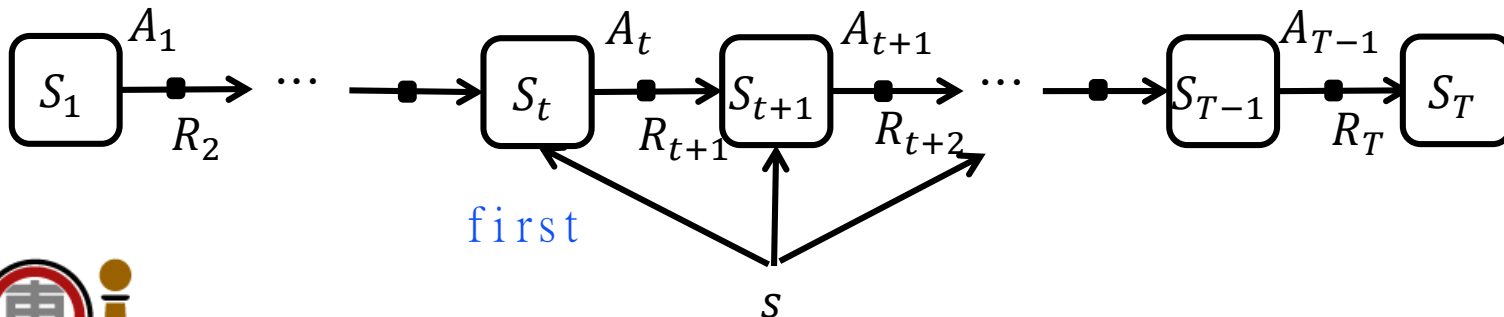


Monte-Carlo Policy Evaluation

- Goal: learn v_π from episodes of experience under policy π
 $S_1, A_1, R_2, \dots, S_k \sim \pi$
- Recall that the return is the total discounted reward:
$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$
- Recall that the value function is the expected return:
$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s]$$
- Monte-Carlo policy evaluation uses empirical mean return instead of expected return

Monte-Carlo Policy Evaluation (cont.)

- To evaluate $v_\pi(s)$ at state s
 - **Every** time-step t that state s is visited in an episode,
 - ▶ Sometimes, we also consider the **first** time-step t .
 - ▶ Both converge quadratically, so this issue is ignored in this course.
 - Increment counter $N(s) \leftarrow N(s) + 1$
 - Increment total return $S(s) \leftarrow S(s) + G_t$
 - Value is estimated by mean return $V(s) \leftarrow S(s)/N(s)$
- By law of large numbers, $V(s) \rightarrow v_\pi(s)$ as $N(s) \rightarrow \infty$



Incremental Mean

The mean μ_1, μ_2, \dots of a sequence x_1, x_2, \dots can be computed incrementally,

$$\begin{aligned}\mu_k &= \frac{1}{k} \sum_{j=1}^k x_j \\ &= \frac{1}{k} (x_k + \sum_{j=1}^{k-1} x_j) \\ &= \frac{1}{k} (x_k + (k-1)\mu_{k-1}) \\ &= \mu_{k-1} + \frac{1}{k} (x_k - \mu_{k-1})\end{aligned}$$

Incremental Monte-Carlo Updates

- Update $V(s)$ incrementally after episode $S_1, A_1, R_2, \dots, S_T$
- For each state S_t with return G_t

$$N(S_t) \leftarrow N(S_t) + 1$$


$$V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t))$$

- In **non-stationary problems**, it can be useful to track a running mean, i.e. **forget old episodes**.

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

Temporal-Difference Learning

- TD methods **learn directly from episodes of experience**
- TD is **model-free**:
 - no knowledge of MDP transitions / rewards
- TD learns from **incomplete episodes**,
 - by bootstrapping
- TD updates a **guess** towards a **guess**



MC vs. TD

- Goal: learn v_π online from experience under policy π
- Incremental every-visit Monte-Carlo
 - Update value $V(S_t)$ toward actual return G_t
$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$
- Simplest temporal-difference learning algorithm: TD(0)
 - Update value $V(S_t)$ toward estimated return $R_{t+1} + \gamma V(S_{t+1})$
$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$
 - $R_{t+1} + \gamma V(S_{t+1})$ is called the **TD target**
 - $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is called the **TD error**

TD vs. MC (I)

- TD can learn **before** knowing the final outcome
 - TD can learn online after every step
 - MC must wait until end of episode before return is known
- TD can learn **without** the final outcome
 - TD can learn from incomplete sequences
 - MC can only learn from complete sequences
 - TD works in continuing (non-terminating) environments
 - MC only works for episodic (terminating) environments

Bias/Variance Trade-Off

- TD target $R_{t+1} + \gamma V(S_{t+1})$ is **biased estimate** of $v_\pi(S_t)$
 - Return $G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$ is **unbiased estimate** of $v_\pi(S_t)$
 - True TD target $R_{t+1} + \gamma v_\pi(S_{t+1})$ is **unbiased estimate** of $v_\pi(S_t)$
- TD target is **much lower variance than the return**:
 - Return depends on **many** random actions, transitions, rewards
 - TD target depends on **only one** random action, transition, reward

MC vs. TD (II)

- MC has high variance, zero bias
 - Good convergence properties (even with function approximation)
 - Not very sensitive to initial value
 - Very simple to understand and use
- TD has low variance, some bias
 - Usually more efficient than MC
 - TD(0) converges to $v_\pi(s)$ (but not always with function approximation)
 - More sensitive to initial value

Batch MC and TD

- MC and TD converge: $V(s) \rightarrow v_\pi(s)$ as experience $\rightarrow \infty$
- But what about batch solution for finite experience?

$$s_1^1, a_1^1, r_2^1, \dots, s_{T_1}^1$$

$$\vdots$$

$$s_1^k, a_1^k, r_2^k, \dots, s_{T_k}^k$$

- e.g. Repeatedly sample episode $k \in [1, K]$
- Apply MC or TD(0) to episode k



AB Example

Two states A, B; no discounting; 8 episodes of experience

A, 0, B, 0

B, 1

B, 1


B, 1

B, 1

B, 1

B, 1

B, 0



What is $V(A)$, $V(B)$?

Both MC and TD will obtain different values!!

Certainty Equivalence

- MC converges to solution with minimum mean-squared error
 - Best fit to the observed returns

$$\sum_{k=1}^K \sum_{t=1}^{T_k} (G_t^k - V(s_t^k))^2$$

- In the AB example, $V(A) = 0$, $V(B) = 0.75$

- TD(0) converges to solution of max likelihood Markov model
 - Solution to the MDP $\langle \mathcal{S}, \mathcal{A}, \hat{\mathcal{P}}, \hat{\mathcal{R}}, \gamma \rangle$ that best fits the data

$$\hat{p}_s^a = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{t=1}^{T_k} 1(s_t^k, a_t^k, s_{t+1}^k = s, a, s')$$

$$\hat{\mathcal{R}}_s^a = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{t=1}^{T_k} 1(s_t^k, a_t^k = s, a) r_t^k$$

- In the AB example, $V(A) = 0.75$, $V(B) = 0.75$




MC vs. TD (III)

- **TD exploits Markov property**
 - Usually **more efficient in Markov environments**
 - ▶ So, TD works well for MDP problems like 2048.
- **MC does **not** exploit Markov property**
 - Usually **more effective in non-Markov environments**
 - ▶ MC works fine for non-MDP too.


Monte-Carlo Backup

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$




Temporal-Difference Backup

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



Dynamic Programming Backup

$$V(S_t) \leftarrow \mathbb{E}_{\pi}[R_{t+1} + \gamma V(S_{t+1})]$$



Bootstrapping and Sampling

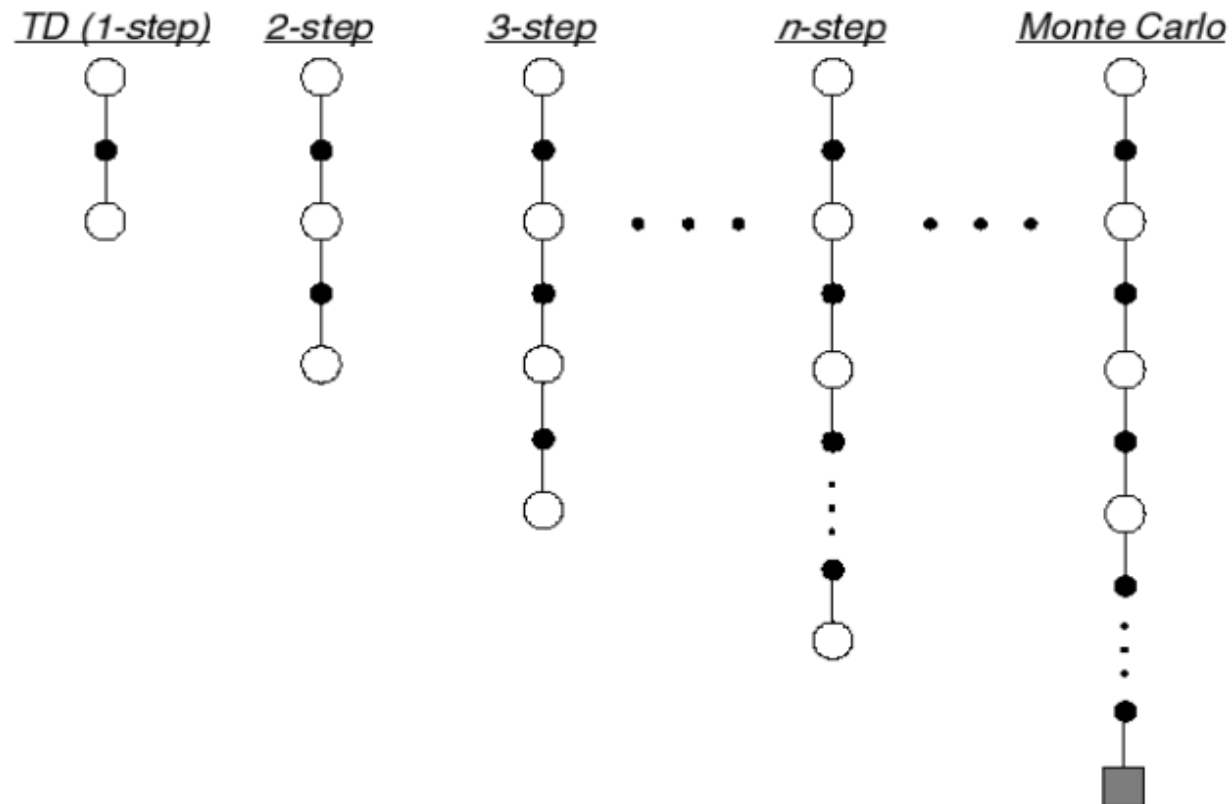
- Bootstrapping: update involves an estimate
 - MC does not bootstrap
 - DP bootstraps
 - TD bootstraps
- Sampling: update samples an expectation
 - MC samples
 - DP does not sample
 - TD samples

General TD Learning

- Review TD
 - Update value $V(S_t)$ toward estimated return $R_{t+1} + \gamma V(S_{t+1})$
$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$
 - $R_{t+1} + \gamma V(S_{t+1})$ is called the **TD target**
- Question: a more general TD target?
- Investigate TD in a more general manner.
- A typical one: TD(λ)

n -Step Prediction

- Let **TD target** look n steps into the future



n -Step Return

- Consider the following n -step returns for $n = 1, 2, \infty$

$$n = 1 \quad G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1})$$

$$n = 2 \quad G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+2})$$

$$\vdots$$


$$n = \infty \quad G_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

- Define the n -step return

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$


- n -step temporal-difference learning

$$V(S_t) \leftarrow V(S_t) + \alpha \left(G_t^{(n)} - V(S_t) \right)$$

 $G_t^{(n)}$


Example of Averaging n -Step Returns

- We can average n -step returns over different n
 - Example:
 - average the 2-step and 4-step returns
- $$\frac{1}{2}G^{(2)} + \frac{1}{2}G^{(4)}$$
- Combines information from two different time-steps
- Next:
 - combine information from all time-steps?



λ -return

- λ -return G_t^λ :
 - combines all n -step returns $G_t^{(n)}$
- Using weight $(1 - \lambda) \lambda^{n-1}$


$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

or (in the case of termination)

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_t^{(n)} + \lambda^{T-t-1} G_t$$


- Forward-view TD(λ)

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$




TD(λ) Weighting Function

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$



Forward-view TD(λ)

- Update value function towards the λ -return
- Forward-view looks into the future to compute G_t^λ
- Like MC, can only be computed from complete episodes



Backward View TD(λ)


- Forward view provides **theory**
- Backward view provides **mechanism**
 - Update online, every step, from incomplete sequences
- Notes:
 - **Consider backward (eligible traces) only when you try to implement it online.** Otherwise, you can ignore it now.

Eligibility Traces

- Credit assignment problem: did bell or light cause shock?
- Frequency heuristic: assign credit to most frequent states
- Recency heuristic: assign credit to most recent states
- Eligibility traces combine both heuristics

$$E_0(s) = 0$$


$$E_t(s) = \gamma \lambda E_{t-1}(s) + 1(S_t = s)$$



Backward View TD(λ)

- Keep an eligibility trace for every state s
- Update value $V(s)$ for every state s
- In proportion to TD-error δ_t and eligibility trace $E_t(s)$

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$
$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$



TD(λ) and TD(0)

- When $\lambda = 0$, only current state is updated

$$E_t(s) = 1(S_t = s)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$

- This is exactly equivalent to TD(0) update

$$V(S_t) \leftarrow V(S_t) + \alpha \delta_t$$



TD(λ) and MC

- When $\lambda = 1$, TD(1) = MC

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_t^{(n)} + \lambda^{T-t-1} G_t = G_t$$

- Consider episodic environments with offline updates
 - ▶ Over the course of an episode, total update for TD(1) is the same as total update for MC
- Consider an episode where s is visited once at time-step k ,

- ▶ TD(1) updates accumulate error online

$$\sum_{t=1}^{T-1} \alpha \delta_t E_t(s) = \alpha \sum_{t=k}^{T-1} \gamma^{t-k} \delta_t = \alpha (G_k - V(S_k))$$

- ▶ By end of episode it accumulates total error

$$\delta_k + \gamma \delta_{k+1} + \gamma^2 \delta_{k+2} + \cdots + \gamma^{T-1-k} \delta_{T-1}$$



Outline

- Model-free prediction: Estimate the value function of an unknown MDP
 - Monte-Carlo (MC) Learning
 - Temporal Difference (TD) Learning
- **Model-free control**: Approximate optimal policies based on the estimation of the value function.
 - On-Policy Monte-Carlo Control
 - On-Policy Temporal-Difference Learning
 - Off-Policy Learning



Model-Free Control

- For most of MDP problems, either:
 - MDP model is **unknown**, but experience can be sampled
 - MDP model is known, but is **too big to use**, **except by samples**
- Model-free control can solve these problems
- On-policy vs. off-policy
 - **On-policy learning**
 - ▶ “Learn on the job”
 - ▶ Learn about policy π from experience sampled from π
 - **Off-policy learning**
 - ▶ “Look over someone's shoulder”
 - ▶ Learn about policy π from experience sampled from π




Generalized Policy Iteration (Recall)

- Policy evaluation \rightarrow Estimate v_π
 - e.g. Iterative policy evaluation
- Policy improvement \rightarrow Generate $\pi' \geq \pi$
 - e.g. Greedy policy improvement



Generalized Policy Iteration With Monte-Carlo Evaluation

- Policy evaluation
 - Monte-Carlo policy evaluation, $V = v_\pi$? Problems: **Model free**?
- Policy improvement
 - Greedy policy improvement?
 - ▶ Problems: **Always choose the same one** (the **best** one)?



Model-Free Policy Iteration Using Action-Value Function

- Greedy policy improvement over $V(s)$ **requires model of MDP**

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} \mathcal{R}_s^a + \mathcal{P}_{ss'}^a V(s')$$


- Greedy policy improvement over $Q(s, a)$ is **model-free**

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a)$$

- Note: this is very important!

Generalized Policy Iteration With Action-Value Function

- Policy evaluation
 - Monte-Carlo policy evaluation, $Q = q_\pi$
- Policy improvement
 - Greedy policy improvement?
 - ▶ Problems: **Always choose the same one** (the **best** one)?



Example of Greedy Action Selection

- There are two doors in front of you,
Always apply the greedy action selection:
 - You open the left door and get reward 0
 $V(left) = 0$
 - You open the right door and get reward +1
 $V(right) = +1$
 - You open the right door and get reward +3
 $V(right) = +2$
 - You open the right door and get reward +2
 $V(right) = +2$
 - \vdots
- Are you sure you've chosen the best door?



ϵ -Greedy Exploration

- ϵ -greedy policy:

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon & \text{if } a^* = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a) \\ \epsilon/m & \text{otherwise} \end{cases}$$

- Exploration

- If you always try the best, you don't explore a real better one.
- With probability ϵ choose an action at random
 - ▶ Simplest idea for ensuring continual exploration
- All m actions are tried with non-zero probability

- Exploitation

- If you always choose at random, you don't exploit the best
- With probability $1 - \epsilon$ choose the greedy action



ϵ -Greedy Policy Improvement

- Theorem

- For any ϵ -greedy policy π , the ϵ -greedy policy π' with respect to q_π is an improvement, $v_{\pi'}(s) \geq v_\pi(s)$

- Proof:


$$\begin{aligned} & q_\pi(s, \pi'(s)) \\ &= \sum_{a \in \mathcal{A}} \pi'(a|s) q_\pi(s, a) \\ &= \epsilon/m \sum_{a \in \mathcal{A}} q_\pi(s, a) + (1 - \epsilon) \max_{a \in \mathcal{A}} q_\pi(s, a) \\ &\geq \epsilon/m \sum_{a \in \mathcal{A}} q_\pi(s, a) + (1 - \epsilon) \sum_{a \in \mathcal{A}} \frac{\pi(a|s) - \epsilon/m}{1 - \epsilon} q_\pi(s, a) \\ &= \sum_{a \in \mathcal{A}} \pi(a|s) q_\pi(s, a) = v_\pi(s) \end{aligned}$$

- Therefore from policy improvement theorem, $v_{\pi'}(s) \geq v_\pi(s)$




Monte-Carlo Policy Iteration

- Policy evaluation
 - Monte-Carlo policy evaluation, $Q = q_\pi$
- Policy improvement
 - ϵ -greedy policy improvement
 - ▶ Converge too, but the proof is not given here.



Monte-Carlo Control

- Policy evaluation
 - Monte-Carlo policy evaluation, $Q \approx q_\pi$
- Policy improvement
 - ϵ -greedy policy improvement
 - ▶ For every episode, improve more slowly by at most a factor of ϵ .




MC vs. TD Control

- Temporal-difference (TD) learning has several advantages over Monte-Carlo (MC)
 - Lower variance
 - Online
 - Incomplete sequences
- Natural idea: use TD instead of MC in our control loop
 - Apply TD to $Q(s, a)$
 - Use ε -greedy policy improvement
 - Update every time-step

Updating Action-Value Functions with Sarsa

$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma Q(S', A') - Q(S, A))$$

- Notice: Interesting naming



Sarsa Algorithm for On-Policy Control

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Repeat (for each step of episode):

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal

- Sarsa converges to the optimal action-value function
- n -step Sarsa – like n -step return
- Sarsa(λ) – like TD(λ)



Off-Policy Learning

- Evaluate **target policy** $\pi(a|s)$ to compute $V_\pi(s)$ or $q_\pi(s, a)$
- While following **behaviour policy** $\mu(a|s)$
 $\{S_1, A_1, R_2, \dots, S_T\} \sim \mu$
- Why is this important?
 - Learn from observing humans or other agents
 - Re-use experience generated from old policies $\pi_1, \pi_2, \dots, \pi_{t-1}$
 - Learn about optimal policy while following exploratory policy
 - Learn about multiple policies while following one policy



Importance Sampling

- Estimate the expectation of a different distribution

$$\begin{aligned} & \mathbb{E}_{X \sim P}[f(X)] \\ &= \sum P(X) f(X) \\ &= \sum Q(X) \frac{P(X)}{Q(X)} f(X) \\ &= \mathbb{E}_{X \sim Q} \left[\frac{P(X)}{Q(X)} f(X) \right] \end{aligned}$$



Importance Sampling for Off-Policy Monte-Carlo

- Use returns generated from μ to evaluate π
- Weight return G_t according to similarity between policies
- Multiply importance sampling corrections along whole episode

$$G_t^{\pi/\mu} = \frac{\pi(A_t|S_t)\pi(A_{t+1}|S_{t+1})}{\mu(A_t|S_t)\mu(A_{t+1}|S_{t+1})} \cdots \frac{\pi(A_T|S_T)}{\mu(A_T|S_T)} G_t$$

- Update value towards corrected return

$$V(S_t) \leftarrow V(S_t) + \alpha \left(G_t^{\pi/\mu} - V(S_t) \right)$$

- Cannot use if μ is zero when π is non-zero
- Importance sampling can dramatically increase variance



Importance Sampling for Off-Policy TD

- Use TD targets generated from μ to evaluate π
- Weight TD target $R + \gamma V(S')$ by importance sampling
- Only need a single importance sampling correction

$$V(S_t) \leftarrow V(S_t) +$$

$$\alpha \left(\frac{\pi(A_t|S_t)}{\mu(A_t|S_t)} (R_{t+1} + \gamma V(S_{t+1})) - V(S_t) \right)$$

- Much lower variance than Monte-Carlo importance sampling (since just one step)
- Policies only need to be similar over a single step



Q-Learning

- We now consider off-policy learning of action-values $Q(s, a)$
- No importance sampling is required
- Next action is chosen using behaviour policy
$$A_{t+1} \sim \mu(\cdot | S_{t+1})$$
- But we consider alternative successor action
$$A' \sim \pi(\cdot | S_{t+1})$$
- And update $Q(S_t, A_t)$ towards value of alternative action
$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A') - Q(S_t, A_t))$$


Off-Policy Control with Q-Learning

- We now allow both behaviour and target policies to improve
- The target policy π is **greedy** w.r.t. $Q(s, a)$
$$\pi(S_{t+1}) = \operatorname{argmax}_{a'} Q(S_{t+1}, a')$$
- The behaviour policy μ is e.g. **ϵ -greedy** w.r.t. $Q(s, a)$
- The Q-learning target then simplifies:

$$\begin{aligned} R_{t+1} + \gamma Q(S_{t+1}, A') \\ &= R_{t+1} + \gamma Q\left(S_{t+1}, \operatorname{argmax}_{a'} Q(S_{t+1}, a')\right) \\ &= R_{t+1} + \max_{a'} \gamma Q(S_{t+1}, a') \end{aligned}$$



Q-Learning Control Algorithm



- $Q(S, A) \leftarrow Q(S, A) + \alpha \left(R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right)$
- Theorem
 - Q-learning control converges to the optimal action-value function, $Q(s, a) \rightarrow q_*(s, a)$

Q-Learning Algorithm for Off-Policy Control

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$;

until S is terminal

