

# Integrating Learning and Planning

I-Chen Wu

- Sutton, R.S. and Barto, A.G., Reinforcement Learning: An Introduction, MIT Press, Cambridge, MA, 1998. (Bible for RL)
  - <http://webdocs.cs.ualberta.ca/~sutton/book/ebook/the-book.html>
  - Chapters 2&9
- David Silver, Online Course for Deep Reinforcement Learning.
  - <http://www.cs.ucl.ac.uk/staff/D.Silver/web/Teaching.html>
  - Chapters 8-9



# Outline

- Introduction
- Model-Based Reinforcement Learning
- Integrated Architectures
- Simulation-Based Search

# Model-Based Reinforcement Learning


- Previous lectures:
  - learn **policy** directly from experience
  - learn **value function** directly from experience
- This lecture:
  - learn **model** directly from experience
  - use **planning** to construct a **value function or policy**
  - **integrate learning and planning** into a single architecture



# Model-Based and Model-Free RL

- **Model-Free RL**
  - No model
  - Learn value function (and/or policy) from experience
- **Model-Based RL**
  - Learn a model from experience
  - Plan value function (and/or policy) from model

# Model-Based RL





# What is a Model?

- A model  $\mathcal{M}$  is a representation of an MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R} \rangle$ , parametrized by  $\eta$
- We will assume state space  $\mathcal{S}$  and action space  $\mathcal{A}$  are known
- So a model  $\mathcal{M} = \langle \mathcal{P}_\eta, \mathcal{R}_\eta \rangle$  represents state transitions  $\mathcal{P}_\eta \approx \mathcal{P}$  and rewards  $\mathcal{R}_\eta \approx \mathcal{R}$

$$S_{t+1} \sim \mathcal{P}_\eta(S_{t+1} | S_t, A_t)$$

$$\mathcal{R}_{t+1} \sim \mathcal{R}_\eta(R_{t+1} | S_t, A_t)$$

- Typically assume conditional independence between state transitions and rewards

$$\mathbb{P}[S_{t+1}, R_{t+1} | S_t, A_t] = \mathbb{P}[S_{t+1} | S_t, A_t] \mathbb{P}[R_{t+1} | S_t, A_t]$$

# Model Learning

- Goal:
  - Estimate model  $\mathcal{M}_\eta$  from experience  $\{S_1, A_1, R_2, \dots, S_T\}$
- This is a supervised learning problem

$$S_1, A_1 \rightarrow R_2, S_2$$

$$S_2, A_2 \rightarrow R_3, S_3$$

$$\vdots$$

$$S_{T-1}, A_{T-1} \rightarrow R_T, S_T$$

- Learning  $s, a \rightarrow r$  is a regression problem
- Learning  $s, a \rightarrow s'$  is a density estimation problem
- Pick loss function, e.g. mean-squared error,  $KL$  divergence, ...
- Find parameters  $\eta$  that minimise empirical loss







# Table Lookup Model

- Model is an explicit MDP,  $\hat{\mathcal{P}}, \hat{\mathcal{R}}$
- Count visits  $N(s, a)$  to each state action pair

$$\hat{\mathcal{P}}_{s,s'}^a = \frac{1}{N(s, a)} \sum_{t=1}^T 1(S_t, A_t, S_{t+1} = s, a, s')$$

$$\hat{\mathcal{R}}_s^a = \frac{1}{N(s, a)} \sum_{t=1}^T 1(S_t, A_t = s, a) R_t$$

- Alternatively
  - At each time-step  $t$ , record experience tuple  $\langle S_t, A_t, R_{t+1}, S_{t+1} \rangle$
  - To sample model, randomly pick tuple matching  $\langle s, a, \cdot, \cdot \rangle$



# AB Example

- Two states  $A, B$ ; no discounting; 8 episodes of experience

$A, 0, B, 0$

$B, 1$

$B, 1$


$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 0$



- We have constructed a table lookup model from the experience



# Planning with a Model

- Given a model  $\mathcal{M}_\eta = \langle \mathcal{P}_\eta, \mathcal{R}_\eta \rangle$
- Solve the MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}_\eta, \mathcal{R}_\eta \rangle$
- Using favorite planning algorithm
  - Value iteration (previous lectures)
  - Policy iteration (previous lectures)
  - Tree search
  - ...

# Sample-Based Planning

- A simple but powerful approach to planning
- Use the model **only to generate samples**
- **Sample** experience from model

$$S_{t+1} \sim \mathcal{P}_\eta(S_{t+1}|S_t, A_t)$$

$$R_{t+1} \sim \mathcal{R}_\eta(R_{t+1}|S_t, A_t)$$

- **Apply model-free RL to samples**, e.g.:
  - Monte-Carlo control
  - Sarsa
  - Q-learning
- **Sample-based planning methods are often more efficient**




# Back to the AB Example

- Construct a table-lookup model from real experience
- Apply model-free RL to sampled experience

Real experience

A, 0, B, 0  
B, 1  
B, 1  
B, 1  
B, 1  
B, 1  
B, 1  
B, 1  
B, 0



Sampled experience

B, 1  
B, 0  
B, 1  
A, 0, B, 1  
B, 1  
A, 0, B, 1  
B, 1  
B, 0

– e.g. Monte-Carlo learning:  $V(A) = 1$ ;  $V(B) = 0.75$



# Planning with an Inaccurate Model

- Given an imperfect model  $\langle \mathcal{P}_\eta, \mathcal{R}_\eta \rangle \neq \langle P, R \rangle$
- Performance of model-based RL is limited to optimal policy for approximate MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}_\eta, \mathcal{R}_\eta \rangle$ 
  - i.e. Model-based RL is only as good as the estimated model
  - When the model is inaccurate, planning process will compute a suboptimal policy
- Solutions
  - when model is wrong, use model-free RL
  - reason explicitly about model uncertainty



# Real and Simulated Experience

- We consider two sources of experience

- Real experience

- Sampled from environment (true MDP)

$$S' \sim \mathcal{P}_{S,S'}^a$$

$$R = \mathcal{R}_S^a$$

- Simulated experience

- Sampled from model (approximate MDP)

$$S' \sim \mathcal{P}_\eta(S' | S, A)$$

$$R = \mathcal{R}_\eta(R | S, A)$$






# Integrating Learning and Planning

- Model-Free RL
  - No model
  - Learn value function (and/or policy) from real experience
- Model-Based RL (using Sample-Based Planning)
  - Learn a model from real experience
  - Plan value function (and/or policy) from simulated experience
- Dyna
  - Learn a model from real experience
  - Learn and plan value function (and/or policy) from real and simulated experience



# Dyna Architecture



# Dyna-Q Algorithm

- Repeat  $n$  times for learning  $Q$  with planning

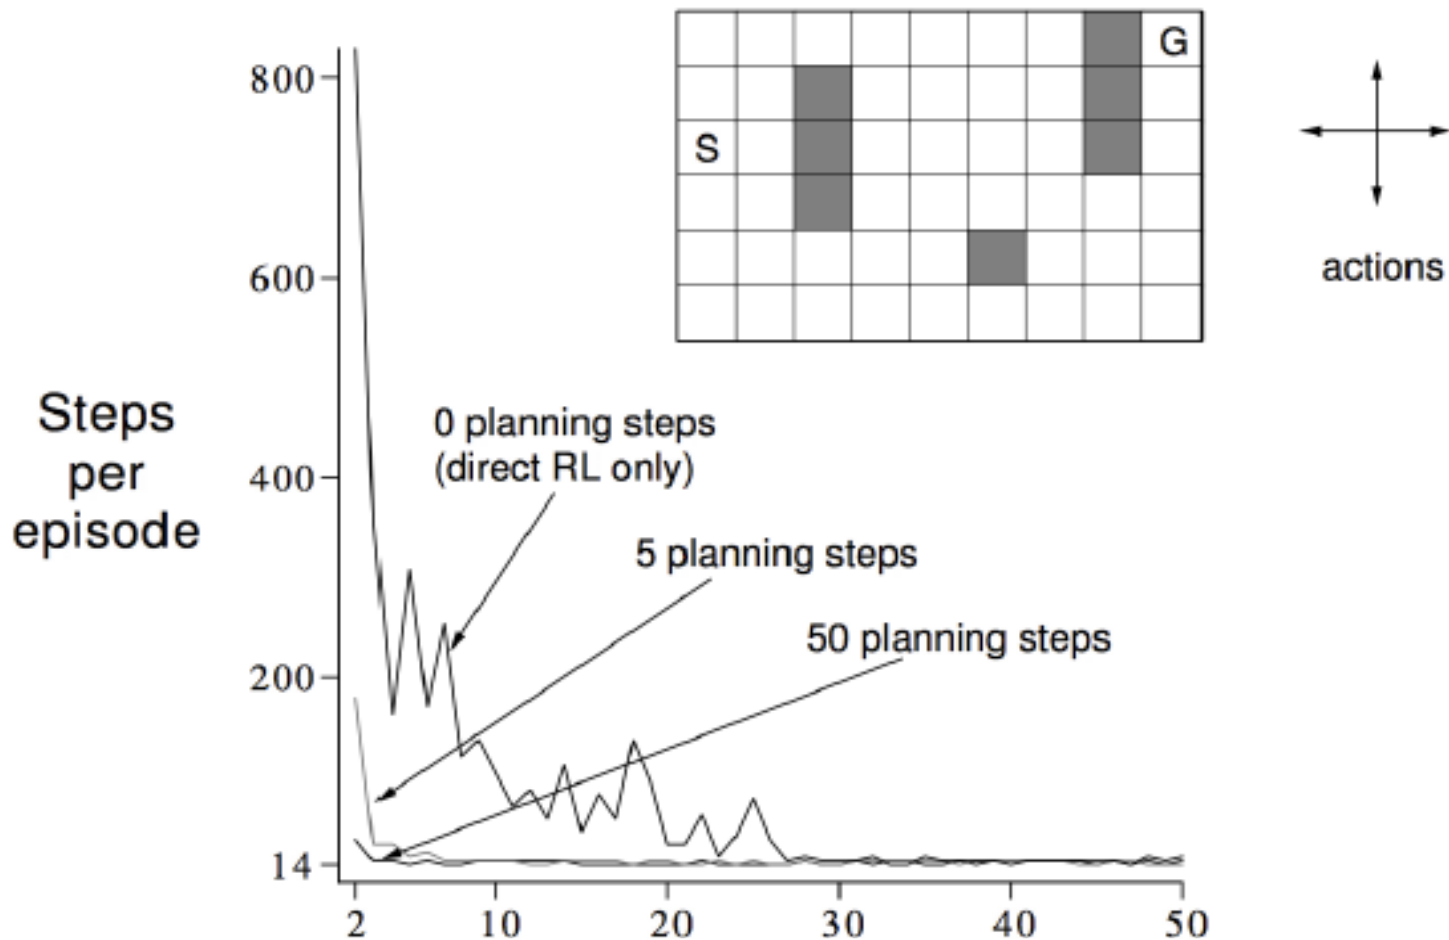
Initialize  $Q(s, a)$  and  $Model(s, a)$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$

Do forever:

- $S \leftarrow$  current (nonterminal) state
- $A \leftarrow \varepsilon$ -greedy( $S, Q$ )
- Execute action  $A$ ; observe resultant reward,  $R$ , and state,  $S'$
- $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$
- $Model(S, A) \leftarrow R, S'$  (assuming deterministic environment)
- Repeat  $n$  times:
  - $S \leftarrow$  random previously observed state
  - $A \leftarrow$  random action previously taken in  $S$
  - $R, S' \leftarrow Model(S, A)$
  - $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$




# Dyna-Q on a Simple Maze




# Dyna-Q with an Inaccurate Model

- When the changed environment is **harder**




# Dyna-Q with an Inaccurate Model (2)

- When the changed environment is **easier**



# Forward Search

- **Forward search** algorithms select the best action by **lookahead**
  - build a search tree with the current state  $S_t$  at the root
  - use a model of the MDP to look ahead




- No need to solve whole MDP, **just sub-MDP starting from now**



# Simulation-Based Search

- Forward search paradigm using sample-based planning
  - Simulate episodes of experience from now **with the model**
  - **Apply model-free RL** to simulated episodes







# Simple Monte-Carlo Search

- Given a model  $\mathcal{M}_v$  and a **simulation policy  $\pi$**
  - For each action  $a \in A$  from current (real) state  $s_t$ 
    - Simulate  $K$  episodes
- $$\{\mathbf{s}_t, \mathbf{a}, R_{t+1}^k, s_{t+1}^k, A_{t+1}^k, \dots, S_T^k\}_{k=1}^K \sim \mathcal{M}_v, \pi$$
- Evaluate actions by mean return (Monte-Carlo evaluation)

$$Q(s_t, a) = \frac{1}{K} \sum_{k=1}^K G_t \xrightarrow{P} q_\pi(s_t, a)$$

- Select current (real) action with maximum value

$$a_t = \operatorname{argmax}_{a \in A} Q(s_t, a)$$



# Monte-Carlo Tree Search (Evaluation)

- Given a model  $M_v$
- Simulate  $K$  episodes from current state  $s_t$  using current simulation policy  $\pi$

$$\{s_t, A_t^k, R_{t+1}^k, s_{t+1}^k, \dots, s_T^k\}_{k=1}^K \sim \mathcal{M}_v, \pi$$

- Build a search tree containing visited states and actions
- Evaluate states  $Q(s, a)$  by mean return of episodes from  $s, a$

$$Q(s, a) = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{u=1}^T 1(S_u, A_u = s, a) G_u \xrightarrow{P} q_\pi(s, a)$$

- After search is finished, select current (real) action with maximum value in search tree

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q(s_t, a)$$



# Monte-Carlo Tree Search (Simulation)

- In MCTS, the simulation policy  $\pi$  improves
- Each simulation consists of two phases (in-tree, out-of-tree)
  - Tree policy (improves): pick actions to maximize  $Q(S, A)$
  - Default policy (fixed): pick actions randomly
- Repeat (each simulation)
  - Evaluate states  $Q(S, A)$  by Monte-Carlo evaluation
  - Improve tree policy, e.g. by  $\epsilon$  – greedy ( $Q$ )
- Notes:
  - Monte-Carlo control applied to simulated experience
  - Converges on the optimal search tree,  $Q(S, A) \rightarrow q_*(S, A)$





## Case Study: the Game of Go

- The ancient oriental game of Go is 2500 years old
- Considered to be the hardest classic board game
- Considered a grand challenge task for AI (John McCarthy)
- Traditional game-tree search has failed in Go



# Rules of Go

- Usually played on 19x19, also 13x13 or 9x9 board
- Simple rules, complex strategy
- Black and white place down stones alternately
- Surrounded stones are captured and removed
- The player with more territory wins the game



# Position Evaluation in Go

- How good is a position  $s$ ?
- Reward function (undiscounted):

$$R_t = 0 \text{ for all non-terminal steps } t < T$$

$$R_T = \begin{cases} 1 & \text{if Black wins} \\ 0 & \text{if White wins} \end{cases}$$

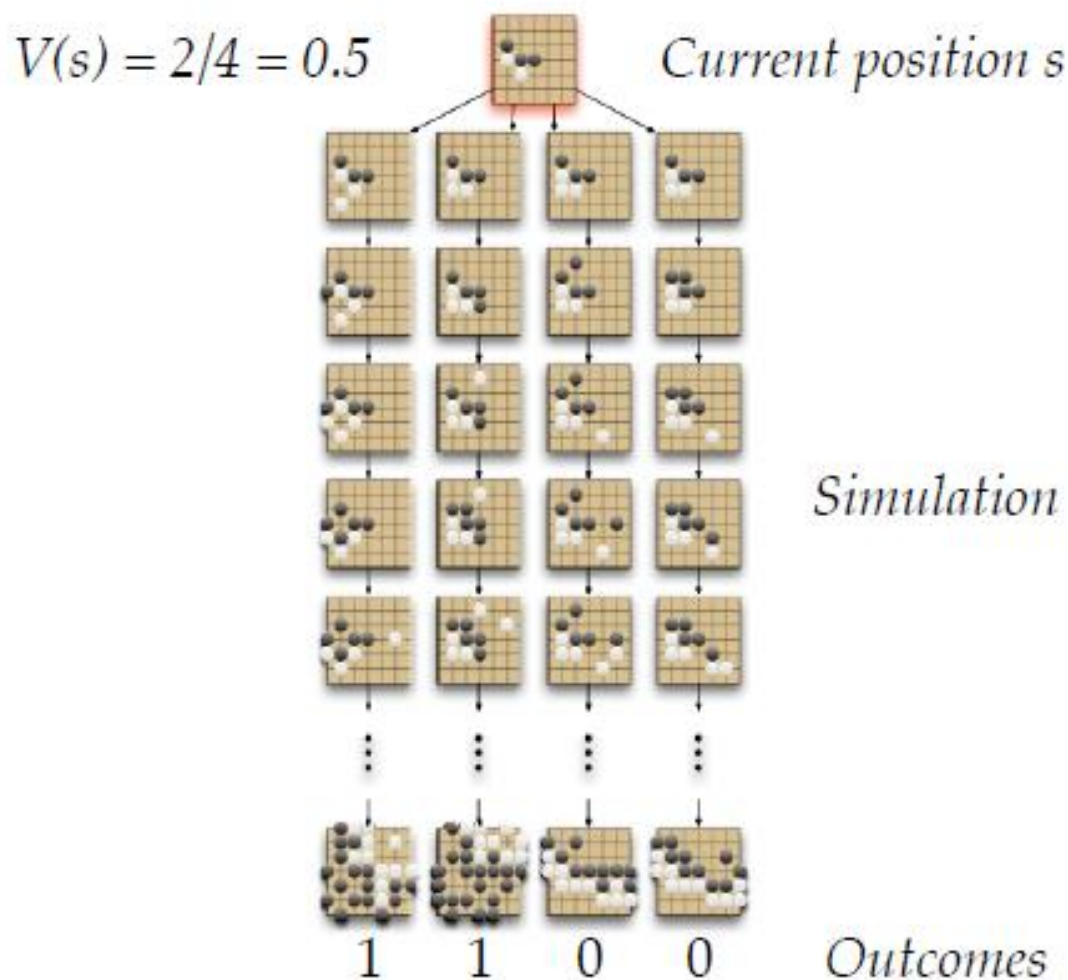
- Policy  $\pi = \langle \pi_B, \pi_W \rangle$  selects moves for both players
- Value function (how good is position  $s$ ):

$$v_\pi(s) = \mathbb{E}_\pi[R_T \mid S = s] = \mathbb{P}[\text{Black wins} \mid S = s]$$

$$v_*(s) = \max_{\pi_B} \min_{\pi_W} v_\pi(s)$$




# Monte-Carlo Evaluation in Go







# Applying Monte-Carlo Tree Search (1)




# Applying Monte-Carlo Tree Search (2)




# Applying Monte-Carlo Tree Search (3)



# Applying Monte-Carlo Tree Search (4)



# Applying Monte-Carlo Tree Search (5)




# Advantages of MC Tree Search

- Highly selective best-first search
- Evaluates states dynamically (unlike e.g. DP)
- Uses sampling to break curse of dimensionality
- Works for “black-box” models (only requires samples)
- Computationally efficient, anytime, parallelizable


# Go – One of the Most Popular Games

- Game tree complexity: about  $10^{360}$ 
  - It is just impossible to try all moves.



# Can Alpha-Beta Search Work for Go?

- Alpha-Beta Search
  - Very successful for many games such as **chess**.
    - ▶ **Almost dominate all computer games before 2006.**
    - ▶ This is what Deep Blue used.
- The key for chess: evaluate position accurately and efficiently.  
E.g., features:
  - King: 1000
  - Queen: 200
  - Rook: 100
  - Knight: 80
  - Bishop: 70
  - Pawn: 30
  - Guarded Pawns: 30
  - Guarded Knights: 40
  - ...
- Problem for chess:
  - need to **consult with experts for feature values.**






# Why not alpha-beta search for Go?

- No simple heuristics like chess.
  - Only black/white pieces (no difference)
- Must know life-and-death
  - But, all are correlated.
    - ▶ Like the lower-right one.
  - But, this is very complex.


Since no simply heuristics to evaluate,

- Why not use Monte-Carlo?
- Calculate it based on stochastics.



# Rules Overview Through a Game (opening 1)


- Black/White move alternately by putting one stone on an intersection of the board.



The example was given by B. Bouzy at CIG'07.

# Rules Overview Through a Game (opening 2)


- Black and White aims at surrounding large « zones »



# Rules Overview Through a Game


(atari 1)

- A white stone is put into « atari » : it has only one liberty left.




# Rules Overview Through a Game (defense)

- White plays to connect the one-liberty stone yielding a four-stone white string with 5 liberties.




# Rules Overview Through a Game (atari 2)

- It is White's turn. One black stone is atari.




## (capture 1)

- White plays on the last liberty of the black stone which is removed



# Rules Overview Through a Game (human end of game)

- The game ends when the two players pass. (Experts would stop here)






# Rules Overview Through a Game


## (contestation 1)

- White contests the black « territory » by playing inside.




# Rules Overview Through a Game (contestation 2)

- White contests black territory, but the 3-stone white string has one liberty left




# Rules Overview Through a Game (follow up 1)

- Black has captured the 3-stone white string



# Rules Overview Through a Game (follow up 2)


- White lacks liberties...



# Rules Overview Through a Game


## (follow up 3)

- Black suppresses the last liberty of the 9-stone string
- Consequently, the white string is removed



# Rules Overview Through a Game (follow up 4)


- Contestation is going on. White has captured four black stones.



# Rules Overview Through a Game



## (concrete end of game)

- The board is covered with either stones or « eyes ».
- Programs know to end.



# Stochastics

- Calculate values based on stochastics.
  - Good example: calculate  $\pi$ .





# Multi-Armed Bandit Problem

## (吃角子老虎問題)

- Assume that you have infinite plays
  - How to choose the one with the maximal average return?



# Exploration vs. Exploitation

## ● Example for the exploration vs exploitation dilemma

- **Exploration:** is a long-term process, with a risky, uncertain outcome.
- **Exploitation:** by contrast is short-term, with immediate, relatively certain benefits



# Deterministic Policy: UCB1

- UCB: Upper Confidence Bounds. [Auer *et al.*, 2002]
- Observed rewards when playing machine  $i$ :  $X_{i,1}, X_{i,2}, \dots$
- Initialization: Play each machine once.

- Loop:
  - Play machine  $j$  that maximizes,  $\bar{X}_j + \sqrt{\frac{2 \log n}{T_j(n)}}$

where  $n$  is the overall number of plays done so far,


$$\bar{X}_{i,s} = \frac{1}{s} \sum_{j=1}^s X_{i,j} \quad , \quad \bar{X}_i = \bar{X}_{i,T_i(n)} \quad ,$$

- Key:
  - Ensure optimal machine is played exponentially more often than any other machine.




# Cumulative Regret

- Assume Machines  $M_1, M_2, M_3, M_4, M_5$ 
  - Win rates: 37%, 42%, 47%, 52%, 57%
  - Trial numbers:  $n_1, n_2, n_3, n_4, n_5$ .




# Monte-Carlo Tree Search

- A kind of planning
- A kind of **Reinforcement learning**



# Strength of Go Program after MCTS


- [Schaeffer et al., 2014]



Strength grew fast, after MCTS.



# Example: MC Tree Search in Computer Go







# MC vs. TD search

- For model-free reinforcement learning, bootstrapping is helpful
  - TD learning reduces variance but increases bias
  - TD learning is usually more efficient than MC
  - TD( $\lambda$ ) can be much more efficient than MC
- For simulation-based search, bootstrapping is also helpful
  - TD search reduces variance but increases bias
  - TD search is usually more efficient than MC search
  - TD( $\lambda$ ) search can be much more efficient than MC search
- Question: can we try TD search for 2048?



# TD Search

- Simulate episodes from the current (real) state  $s_t$
- Estimate action-value function  $Q(s, a)$
- For each step of simulation, update action-values by Sarsa
$$\Delta Q(S, A) = \alpha(R + \gamma Q(S', A') - Q(S, A))$$
- Select actions based on action-values  $Q(s, a)$ 
  - e.g.  $\epsilon$ -greedy
- May also use function approximation for  $Q$



# Dyna-2

- In Dyna-2, the agent stores two sets of feature weights
  - Long-term memory
  - Short-term (working) memory
- Long-term memory is updated from real experience using TD learning
  - General domain knowledge that applies to any episode
- Short-term memory is updated from simulated experience using TD search
  - Specific local knowledge about the current situation
- Over value function is sum of long and short-term memories



## Results of TD search in Go

