

**Deep Learning and Practice**  
**Lab 9: Deep Q Network**  
李韓  
0556157  
fm.bigballon@gmail.com

## Episode rewards

---

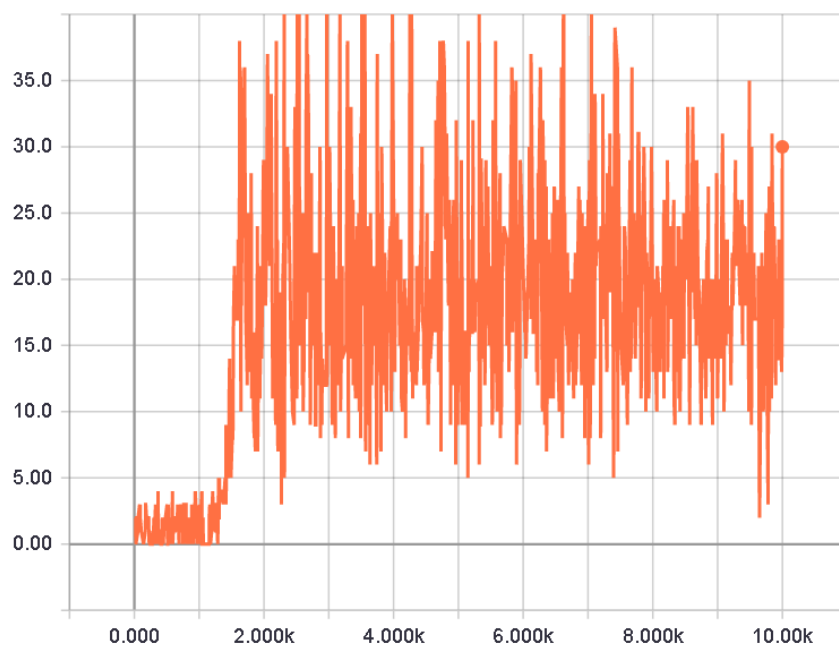
It cost about 30 hours to train the model.

The first time the highest episode reward during training is 78, but I forget to save the plot and logs.

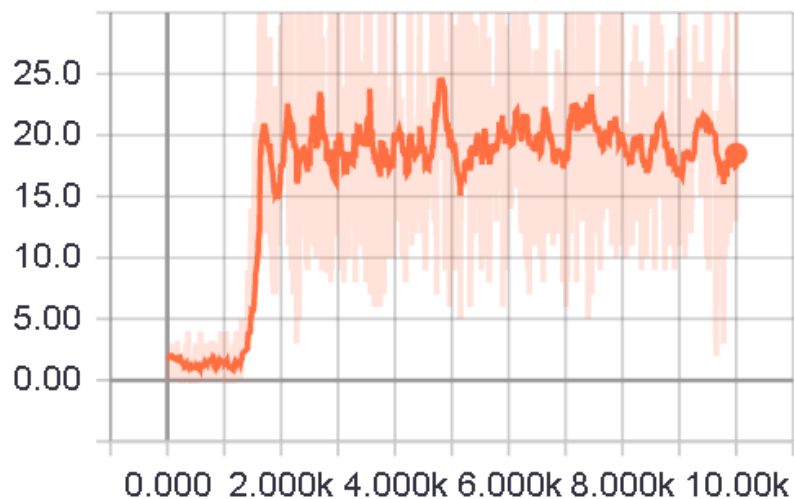
The following figures is the second time's plot:

The highest episode reward during training is 60.

And the average episode reward is about 23.



Episode Reward



# Deep Q network implementation

## ■ Network Structure

According to the Implementation Details, we need to set up 3 convolution layers.

- conv1
- conv2
- conv3

Then we need two fully connected layers

- fc1
- predictions

```
def _build_network(self):
    self.X_pl = tf.placeholder(shape=[None, 84, 84, 4], dtype=tf.uint8, name="X")
    self.y_pl = tf.placeholder(shape=[None], dtype=tf.float32, name="y")
    self.actions_pl = tf.placeholder(shape=[None], dtype=tf.int32, name="actions")

    X = tf.to_float(self.X_pl) / 255.0

    conv1 = tf.contrib.layers.conv2d(X, 32, 8, 4, activation_fn=tf.nn.relu)
    conv2 = tf.contrib.layers.conv2d(conv1, 64, 4, 2, activation_fn=tf.nn.relu)
    conv3 = tf.contrib.layers.conv2d(conv2, 64, 3, 1, activation_fn=tf.nn.relu)

    # fully connected layers
    flattened = tf.contrib.layers.flatten(conv3)
    fc1 = tf.contrib.layers.fully_connected(flattened, 512)
    self.predictions = tf.contrib.layers.fully_connected(fc1, len(VALID_ACTIONS))

    # Get the predictions for the chosen actions only
    batch_size = tf.shape(self.X_pl)[0]
    gather_indices = tf.range(batch_size) * tf.shape(self.predictions)[1] + self.actions_pl
    self.action_predictions = tf.gather(tf.reshape(self.predictions, [-1]), gather_indices)

    # Calculate the loss
    self.losses = tf.squared_difference(self.y_pl, self.action_predictions)
    self.loss = tf.reduce_mean(self.losses)

    # Optimizer Parameters from original paper
    self.optimizer = tf.train.RMSPropOptimizer(0.00025, 0.99, 0.0, 1e-6)
    self.train_op = self.optimizer.minimize(self.loss, global_step=tf.contrib.framework.get_or_create_global_step())

    # Summaries for Tensorboard
    self.summaries = tf.summary.merge([
        tf.summary.scalar("loss", self.loss),
        tf.summary.histogram("loss_hist", self.losses),
        tf.summary.histogram("q_values_hist", self.predictions),
        tf.summary.scalar("max_q_value", tf.reduce_max(self.predictions))
    ])
}
```

## ■ Loss function

1. calculate the [action\_predictions]
2. then calculate the squared\_difference [losses] of (y and action\_predictions)
3. calculate the average of losses [loss]
4. using RMSPropOptimizer to minimize

## ■ Implement update\_target\_network()

```
def update_target_network(sess, behavior_Q, target_Q):  
    e1_params = [t for t in tf.trainable_variables() if t.name.startswith(behavior_Q.scope)]  
    e1_params = sorted(e1_params, key=lambda v: v.name)  
    e2_params = [t for t in tf.trainable_variables() if t.name.startswith(target_Q.scope)]  
    e2_params = sorted(e2_params, key=lambda v: v.name)  
  
    update_ops = []  
    for e1_v, e2_v in zip(e1_params, e2_params):  
        op = e2_v.assign(e1_v)  
        update_ops.append(op)  
  
    sess.run(update_ops)
```

1. using [tf.trainable\_variables] to returns all variables created with trainable=True,
2. then using a list update\_ops to save all assign operators,
3. let sess run the update\_ops to finish assigned between behavior\_Q and target\_Q

## ■ Explain how you implement the training process of deep Q learning

### ■ Populate replay memory

```
action_probs = policy(sess, state, epsilons[min(total_tmp, EXPLORE_STPES-1)])  
action = np.random.choice(np.arange(len(action_probs)), p=action_probs)
```

Using linspace to generate anneal epsilon,

Using epsilon\_greedy\_policy to select action,

Playing the game until replay\_memory less than INIT\_REPLAY\_MEMORY\_SIZE

### ■ Select actions

With probability  $\epsilon$  select a random action  $a_t$

otherwise select  $a_t = \operatorname{argmax}_a Q(\varphi(s_t), a; \theta)$

### ■ Update Epsilon

```
epsilons = np.linspace(INITIAL_EPSILON, FINAL_EPSILON, EXPLORE_STPES)
```

When epsilon increase ,

```
epsilon = epsilons[min(total_t, EXPLORE_STPES-1)]
```

## ■ Prepare minibatch for network update

```
samples = random.sample(replay_memory, BATCH_SIZE)
states_batch, action_batch, reward_batch, next_states_batch, done_batch = map(np.array, zip(*samples))
```

Using random.sample to sample a minibatch from the replay memory (BATCH\_SIZE = 32)

Using map to mapping samples to all informations(states actions rewards next\_states and done)

## Performance – Highest episode reward during training

---

Training 1: the highest episode reward during training is 78.

Training 2: the highest episode reward during training is 60.

